

Online Learning via Convex Geometry, with Applications to Pricing

Adrian Vladu
MIT

Dynamic Pricing Problem





? °



? °



? °



\$0.5



\$2.0



\$1.0





? °



? °



? °

	\$0.5
	\$2.0
	\$1.0



day 1

	x 1
	x 3
	x 0
	?





? °



? °



? °

	\$0.5
	\$2.0
	\$1.0



day 1

	x 1
	x 3
	x 0
	?



\$7





	\$0.5
	\$2.0
	\$1.0



day 1

	x 1
	x 3
	x 0
	?



\$7



X



	\$0.5
	\$2.0
	\$1.0



Regret=6.5

day 1

	x 1
	x 3
	x 0
	?



\$7



X



\$0.5



\$2.0



\$1.0



Regret=6.5



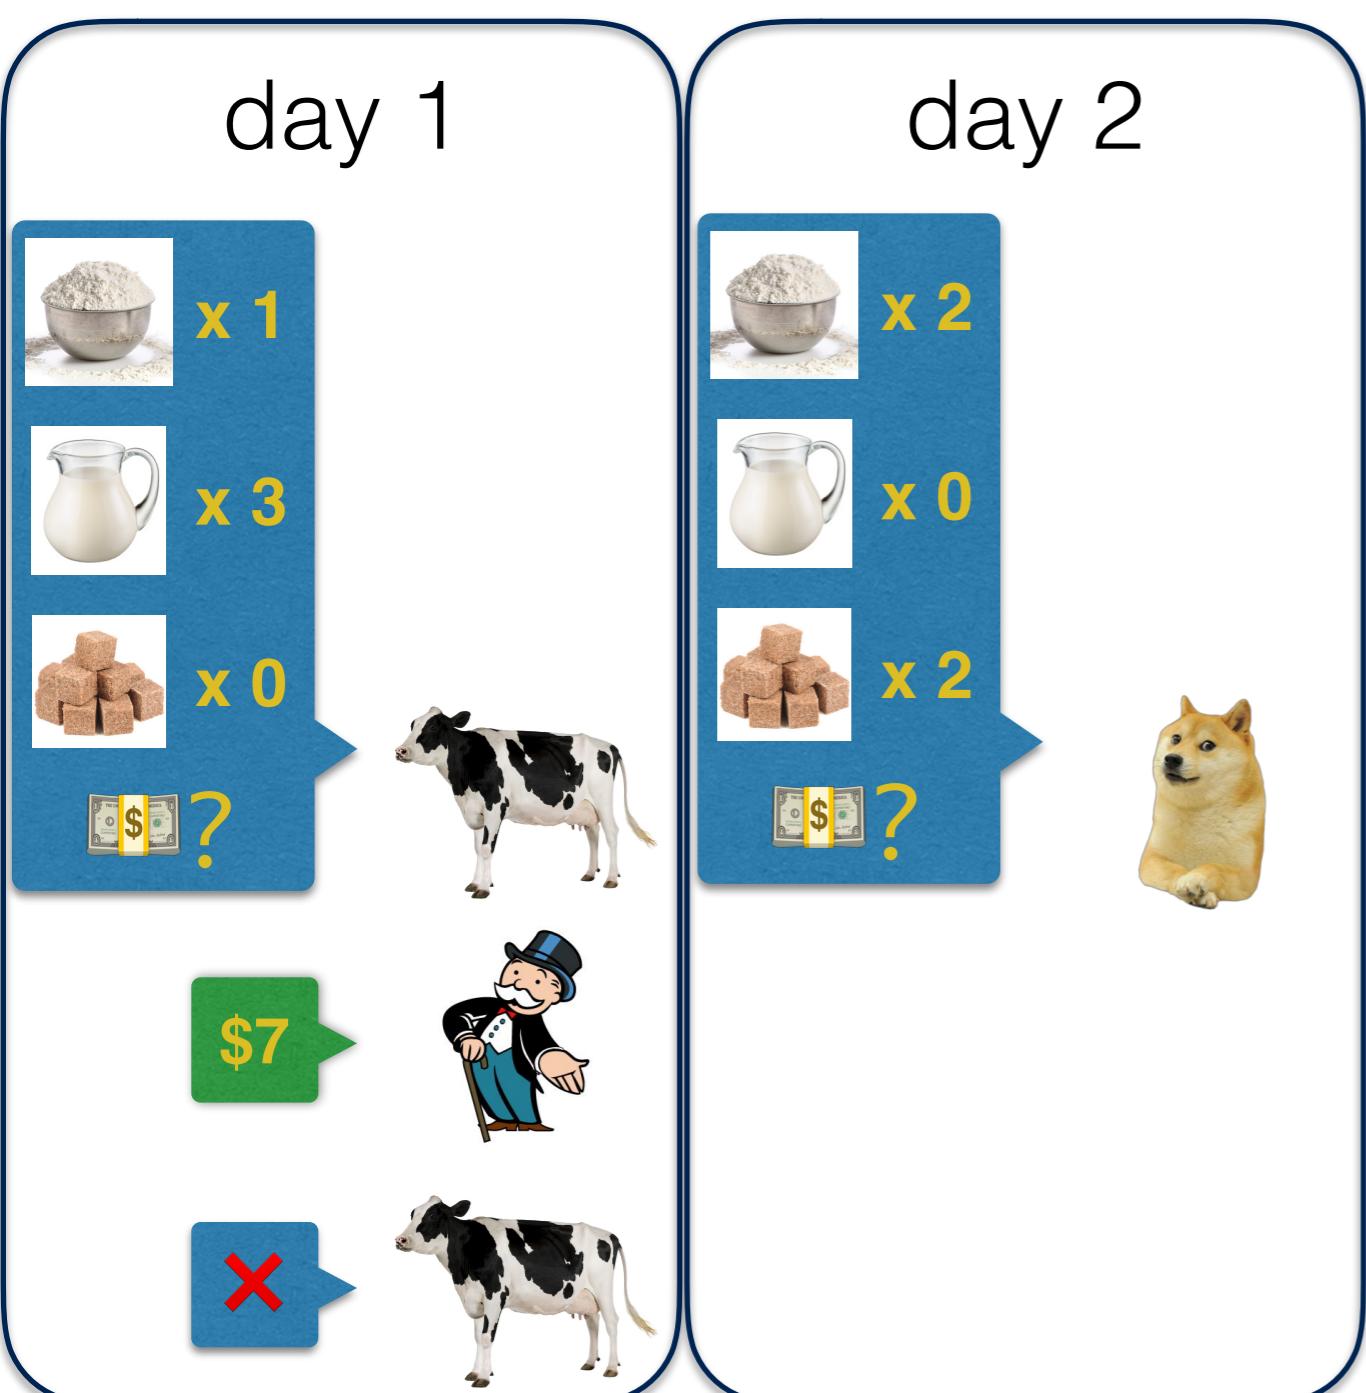
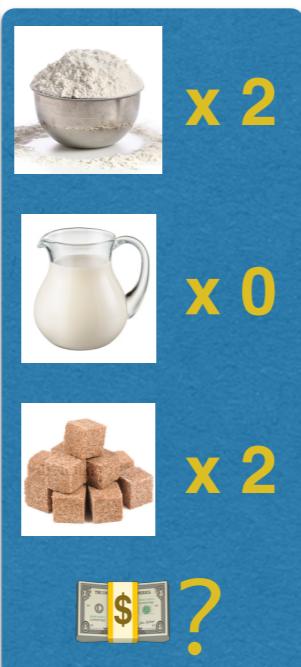
day 1



\$7



day 2





\$0.5



\$2.0



\$1.0



Regret=6.5



day 1



\$7



X



day 2



\$2





\$0.5



\$2.0



\$1.0



Regret=7.5



day 1



\$7



X



day 2



\$2



✓





\$0.5



\$2.0



\$1.0



Regret=7.5



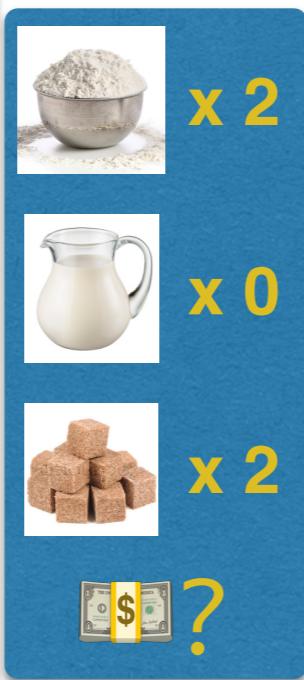
day 1



\$7



X



\$2

✓





\$0.5



\$2.0



\$1.0



Regret=7.5



day 1



x 1



x 3



x 0



?



\$7



X



day 2



x 2



x 0



x 2



?



\$2



day 3



x 0



x 1



x 1



?



\$4





\$0.5



\$2.0



\$1.0



Regret=10.5



day 1



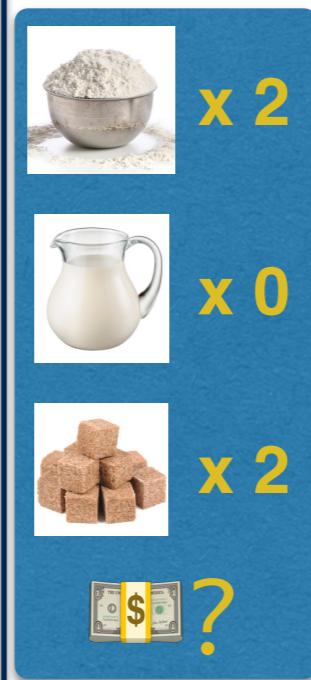
\$7



X



day 2



\$2

✓



day 3



\$4





\$0.5



\$2.0



\$1.0



Regret=10.5

day 1



x 1



x 3



x 0



?



\$7



X



day 2



x 2



x 0



x 2



?



\$2

✓



day 3



x 0



x 1



x 1



?



\$4



day 4



x 1



x 1



x 1



?





	\$0.5
	\$2.0
	\$1.0



Regret=10.5

day 1



\$7



day 2



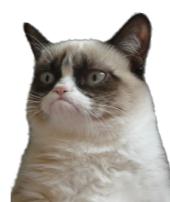
\$2



day 3



\$4



day 4



\$3





	\$0.5
	\$2.0
	\$1.0



Regret=11

day 1



\$7



day 2



\$2



day 3



\$4



day 4



\$3



Dynamic Pricing Problem

- Warm-up: 1-dimensional case

Dynamic Pricing Problem

- Warm-up: 1-dimensional case



Dynamic Pricing Problem

- Warm-up: 1-dimensional case



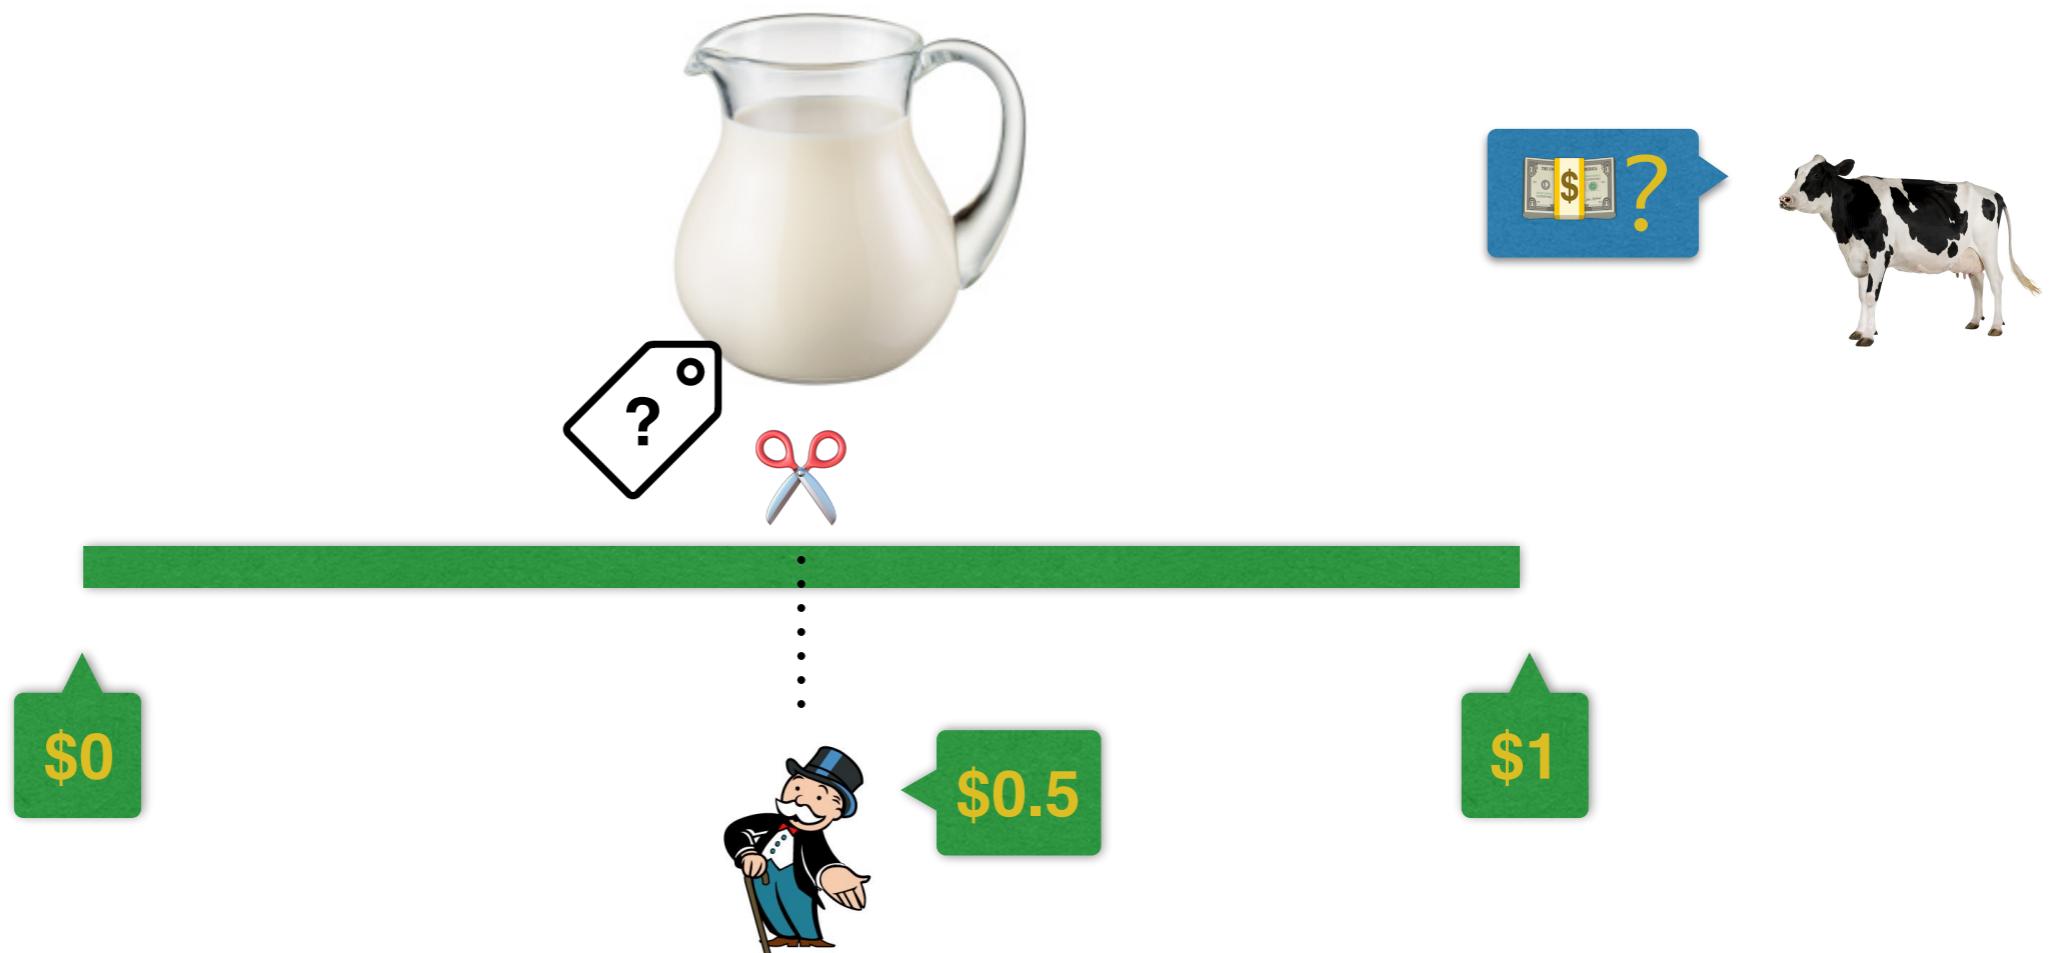
Dynamic Pricing Problem

- Warm-up: 1-dimensional case



Dynamic Pricing Problem

- Warm-up: 1-dimensional case



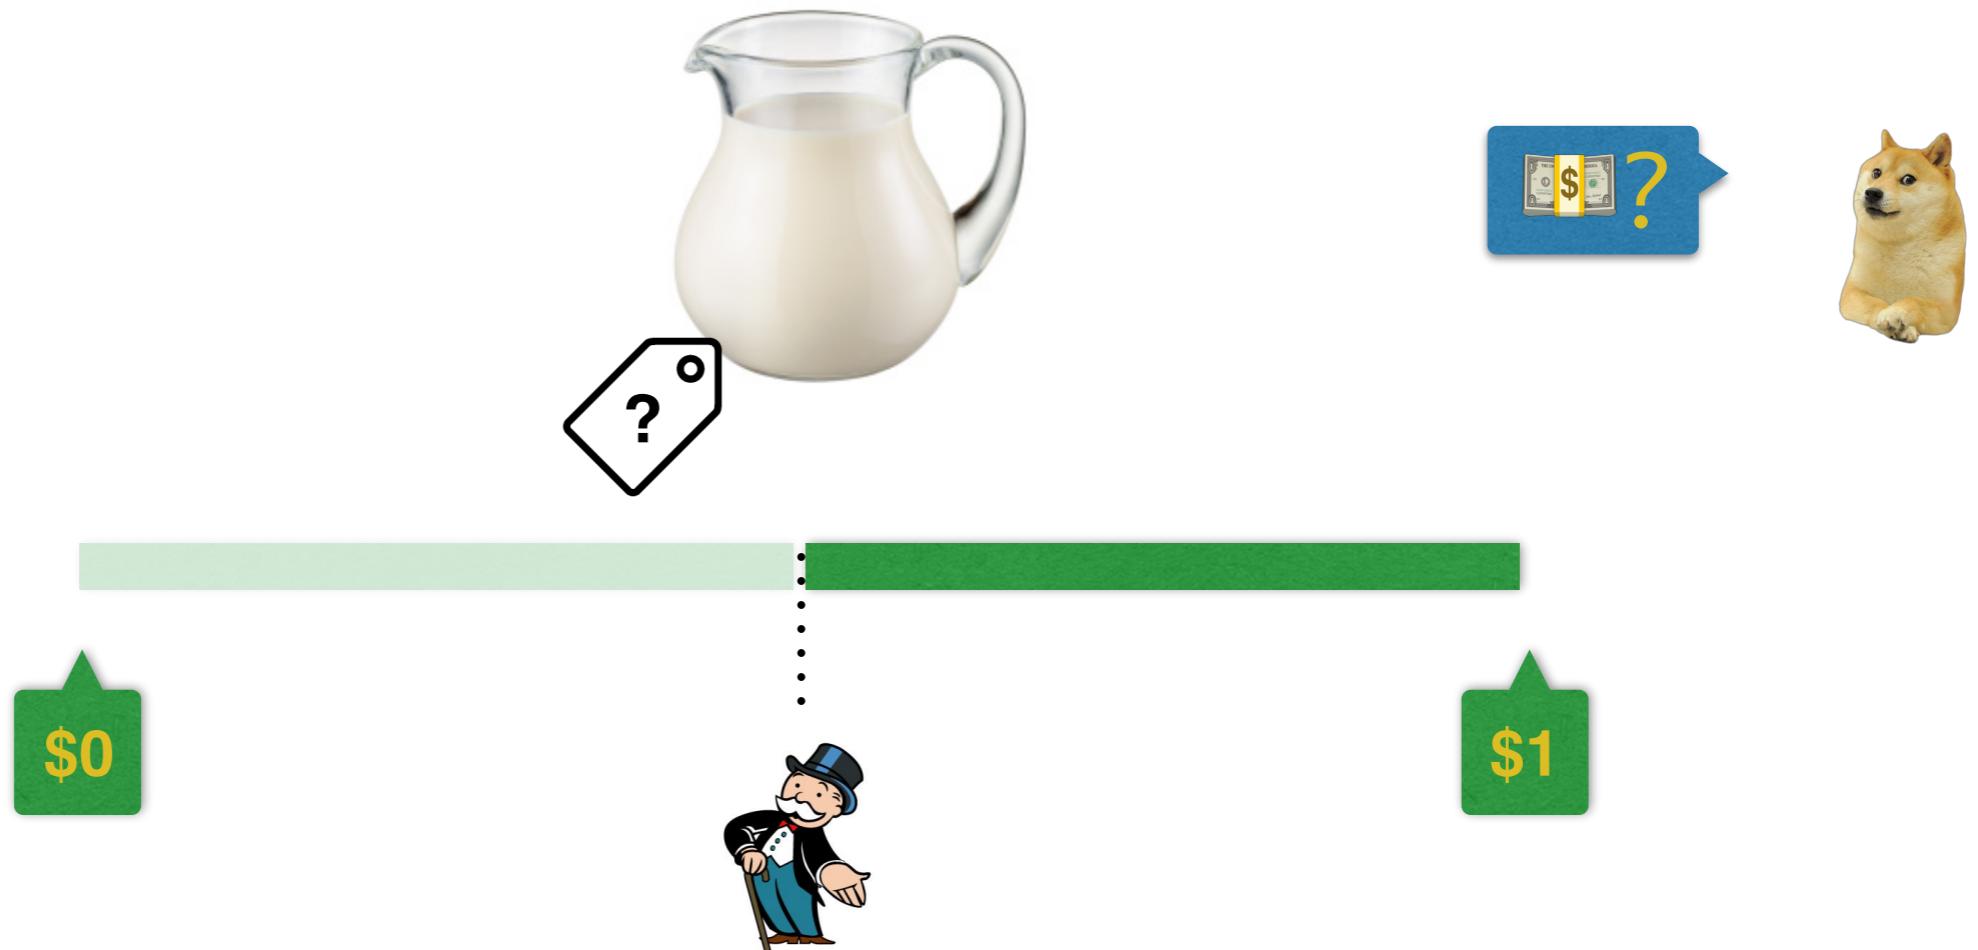
Dynamic Pricing Problem

- Warm-up: 1-dimensional case



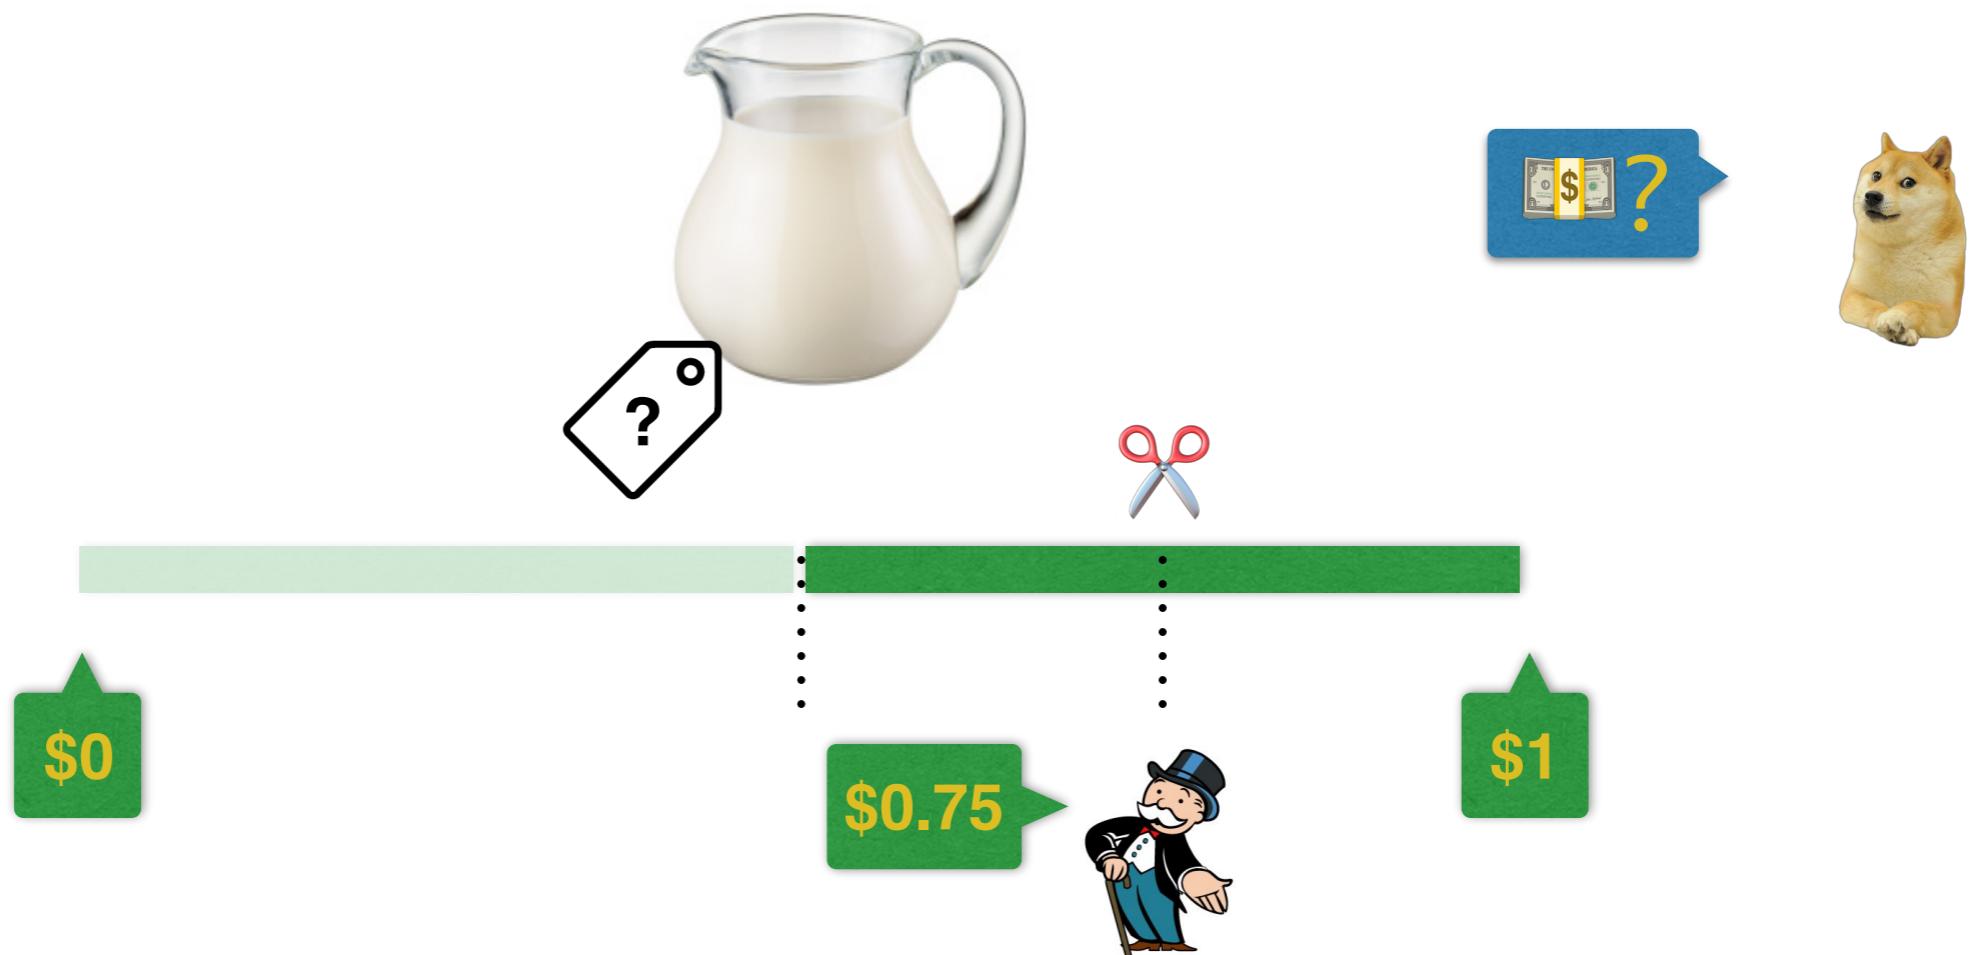
Dynamic Pricing Problem

- Warm-up: 1-dimensional case



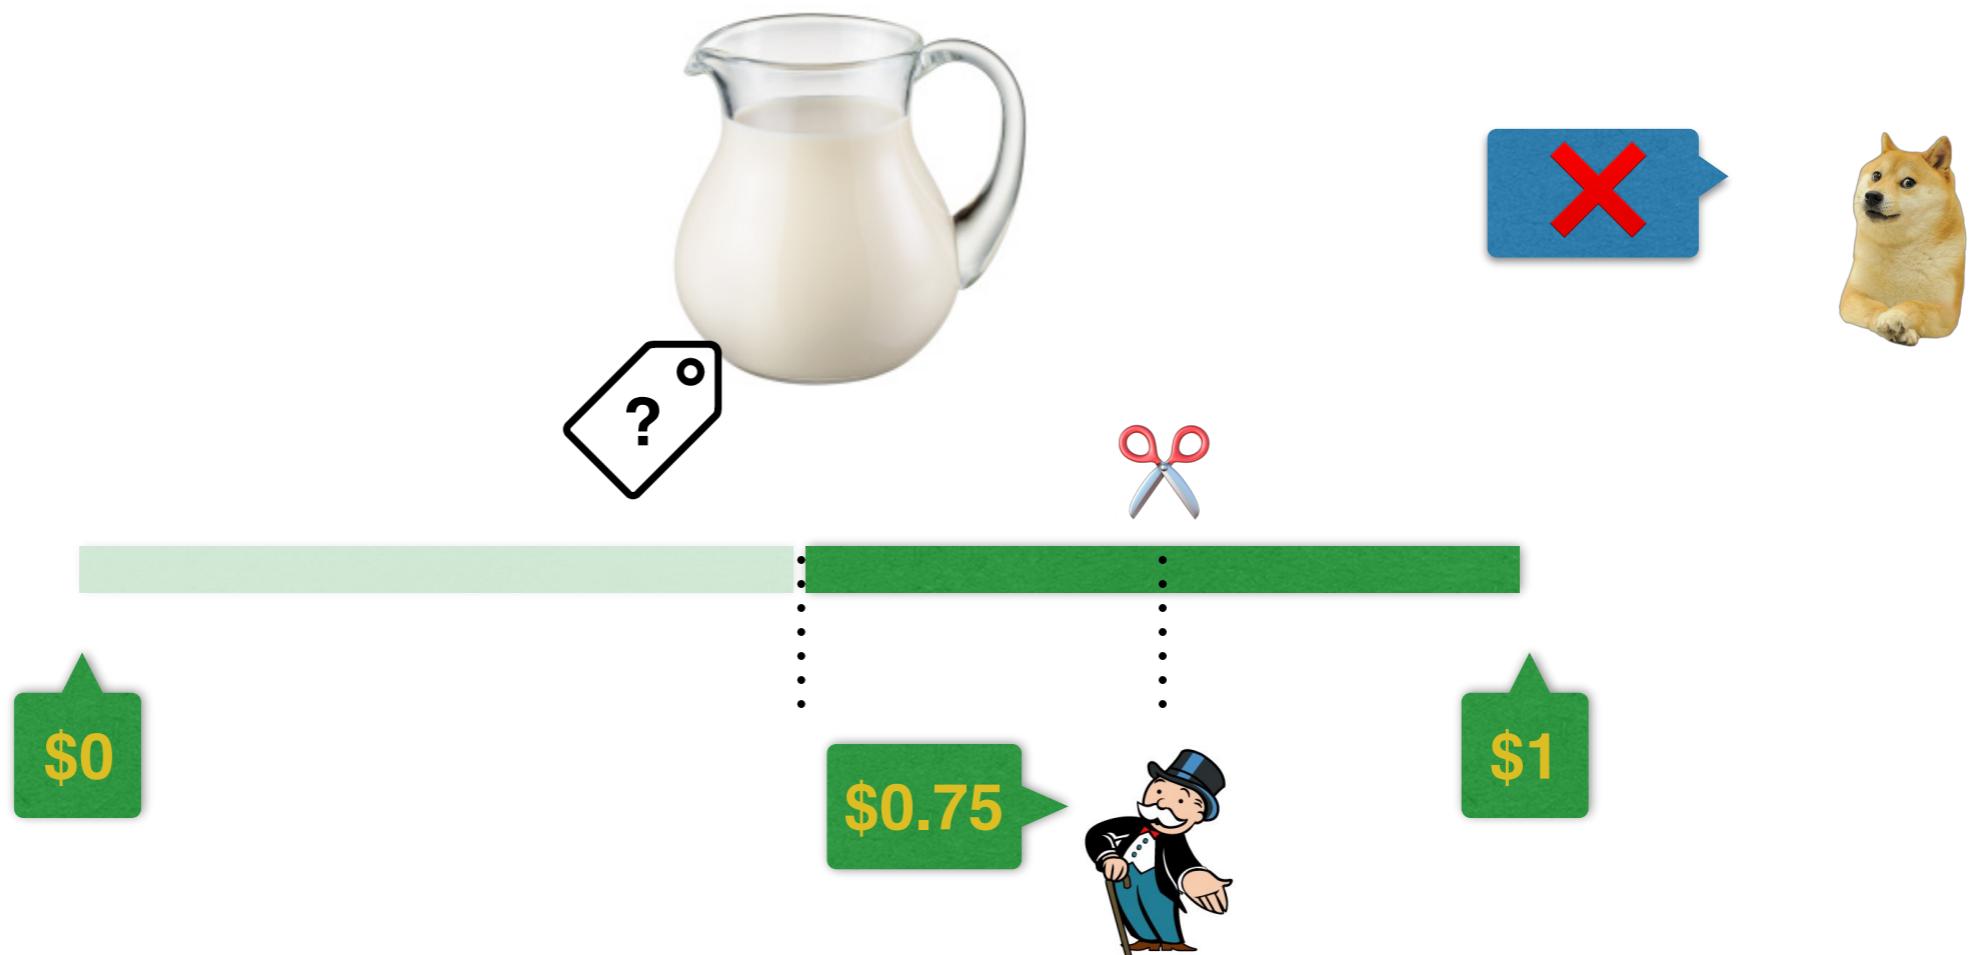
Dynamic Pricing Problem

- Warm-up: 1-dimensional case



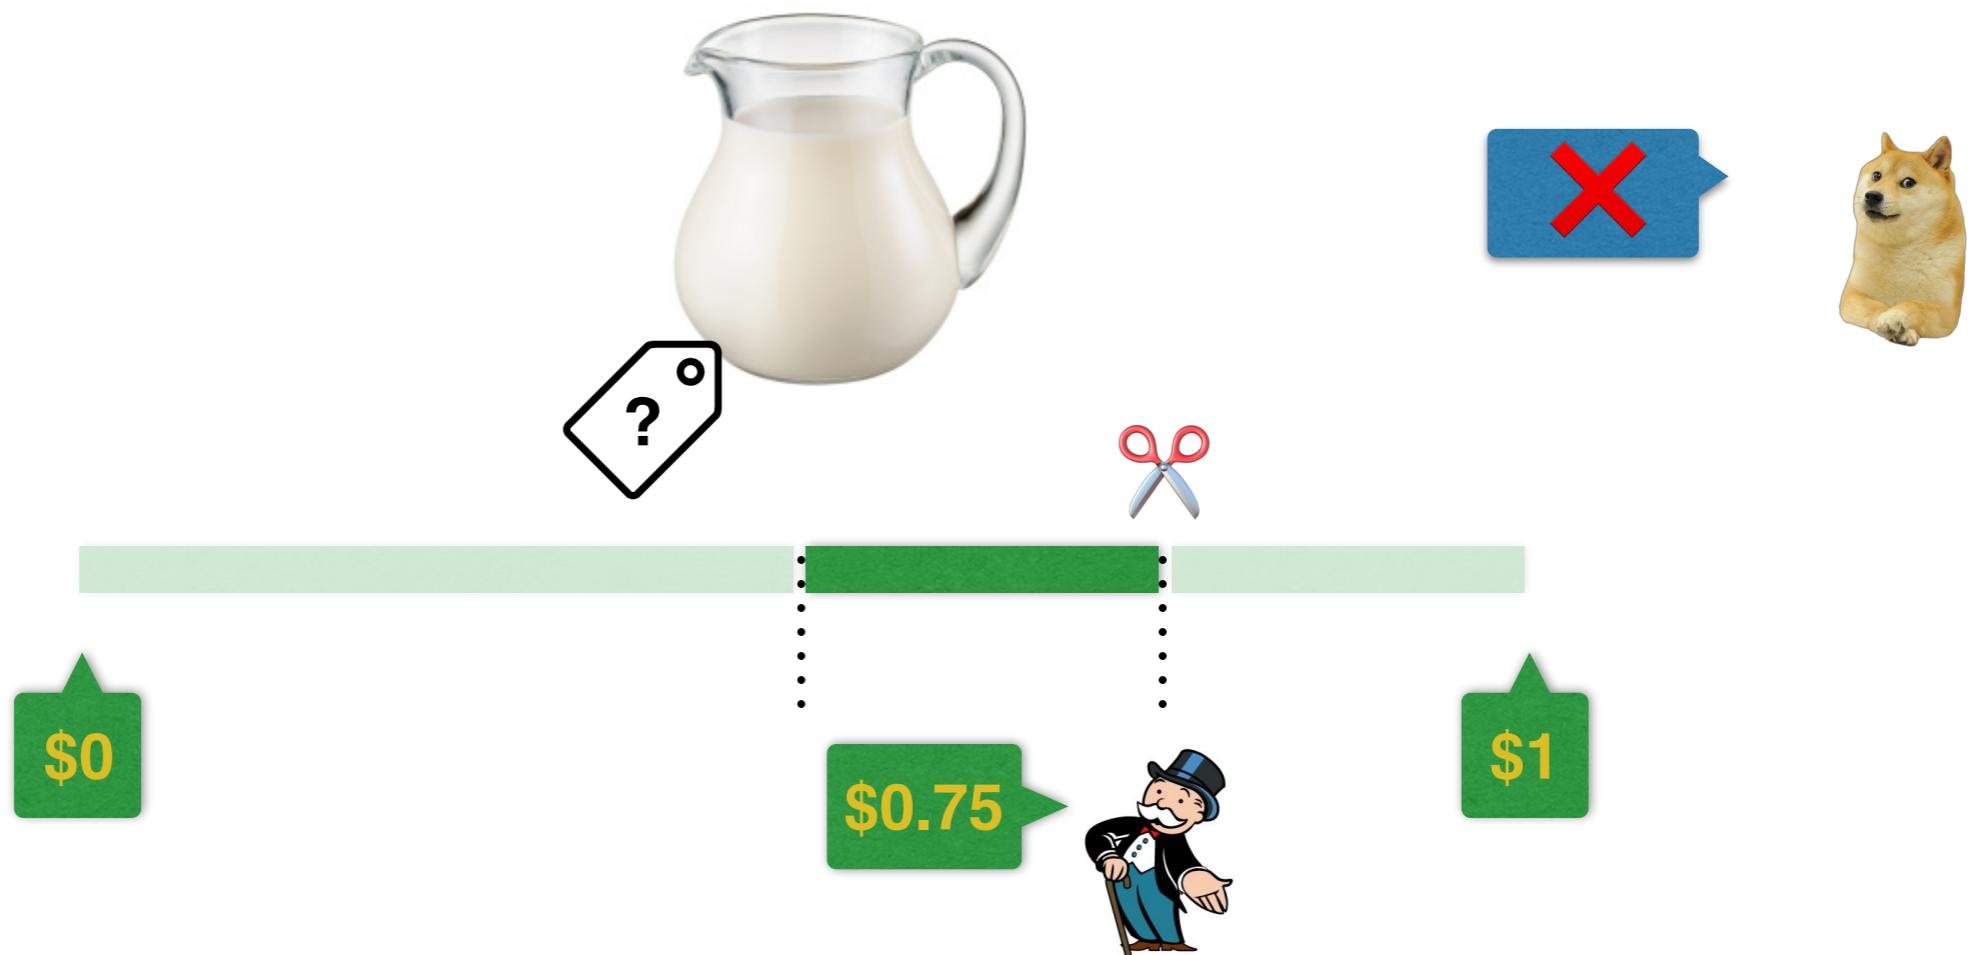
Dynamic Pricing Problem

- Warm-up: 1-dimensional case



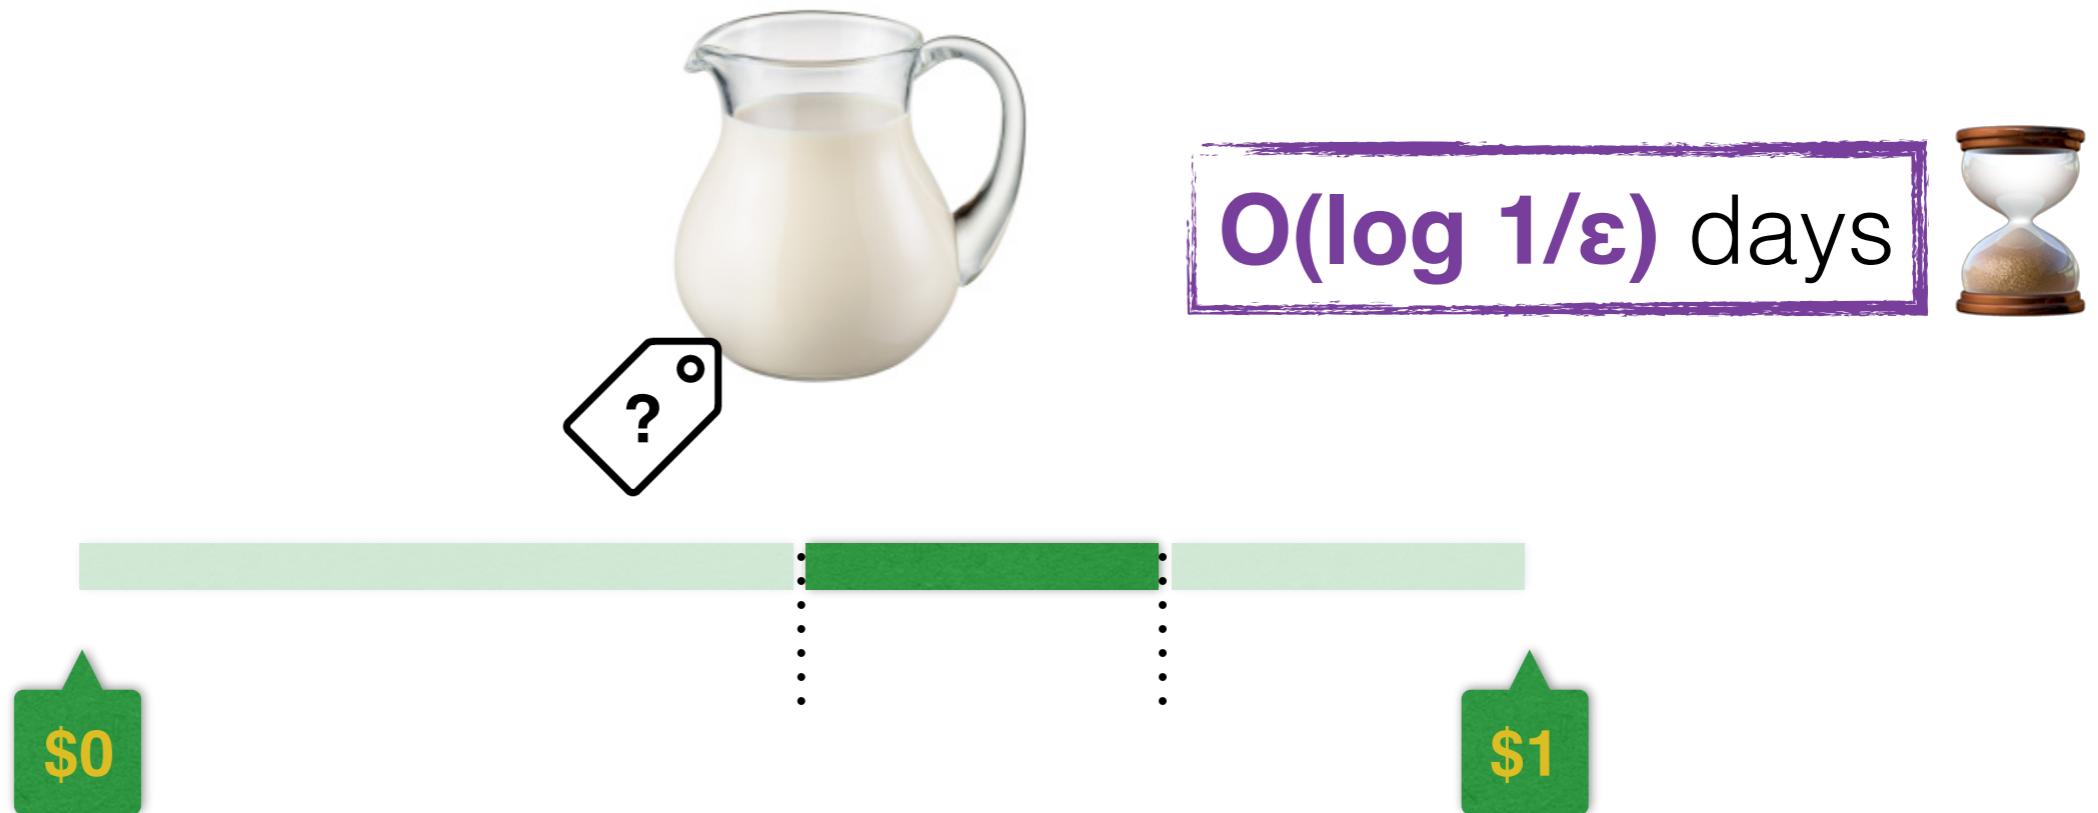
Dynamic Pricing Problem

- Warm-up: 1-dimensional case



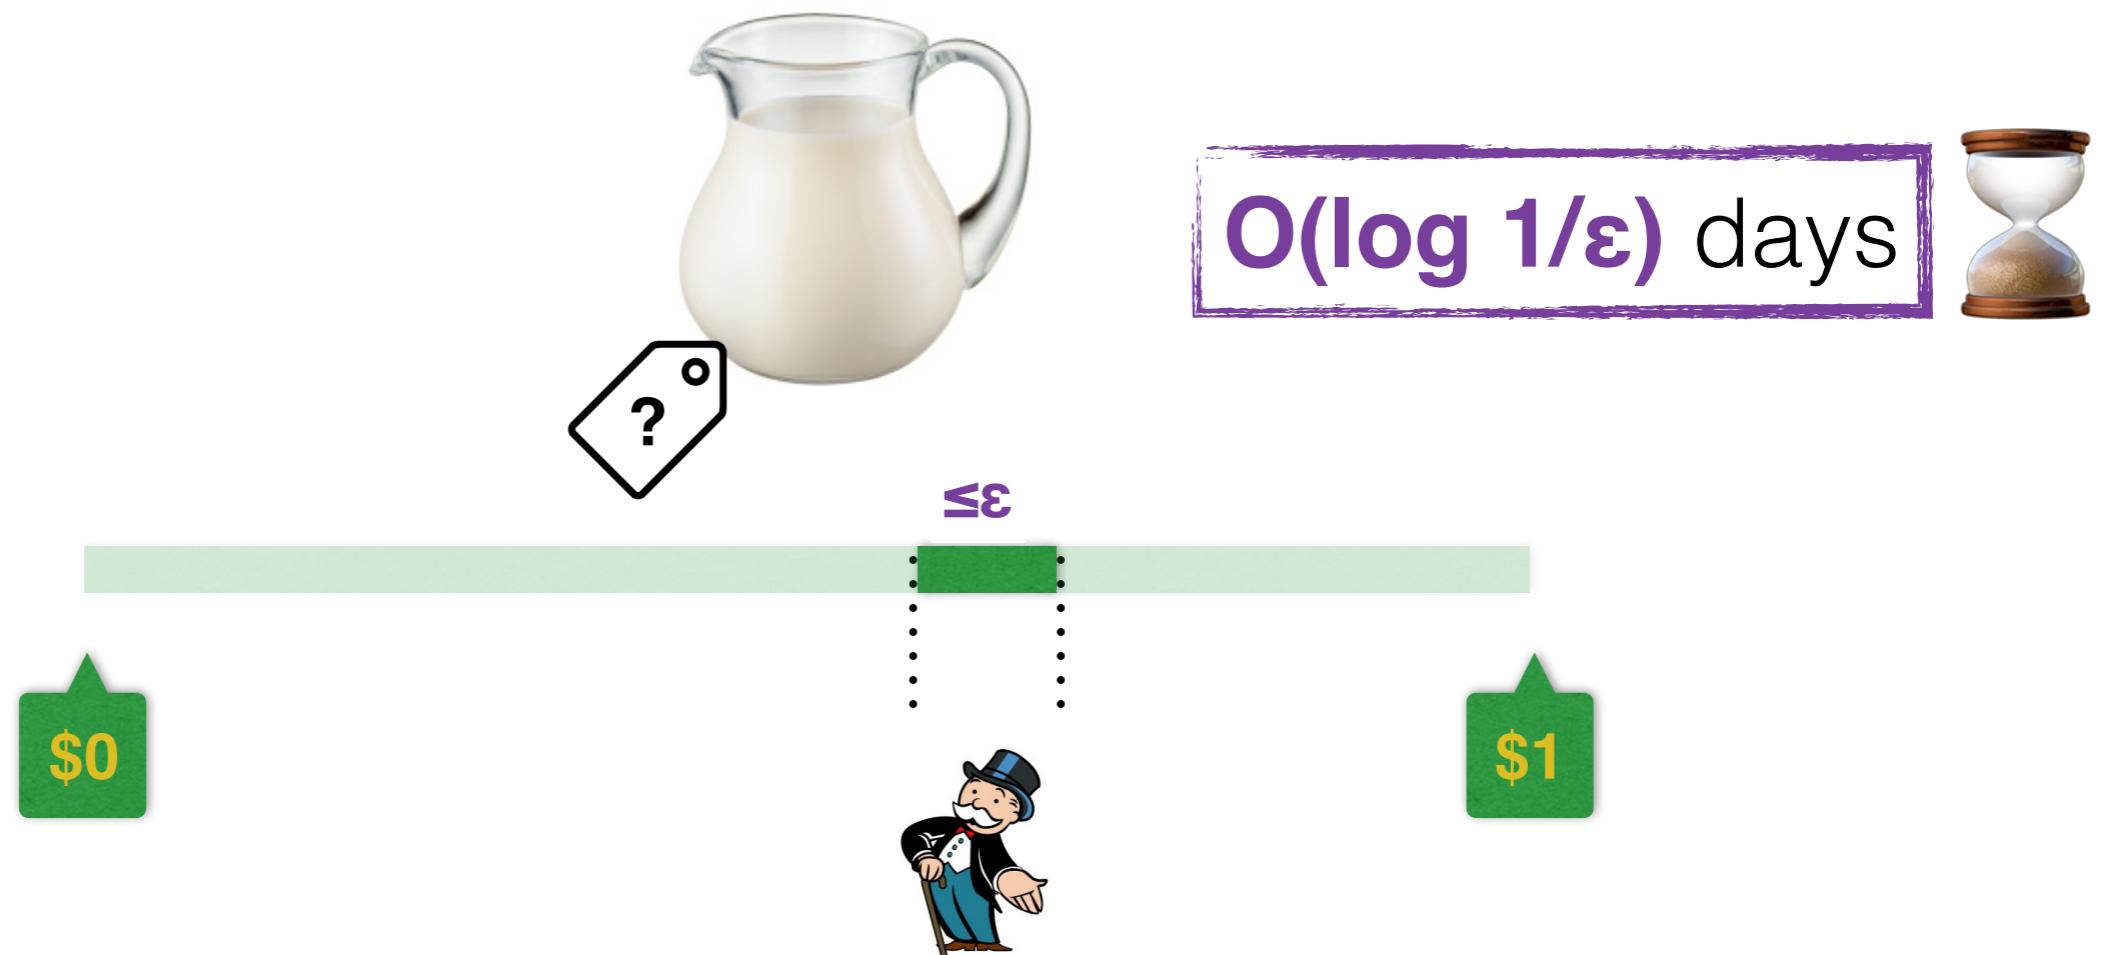
Dynamic Pricing Problem

- Warm-up: 1-dimensional case



Dynamic Pricing Problem

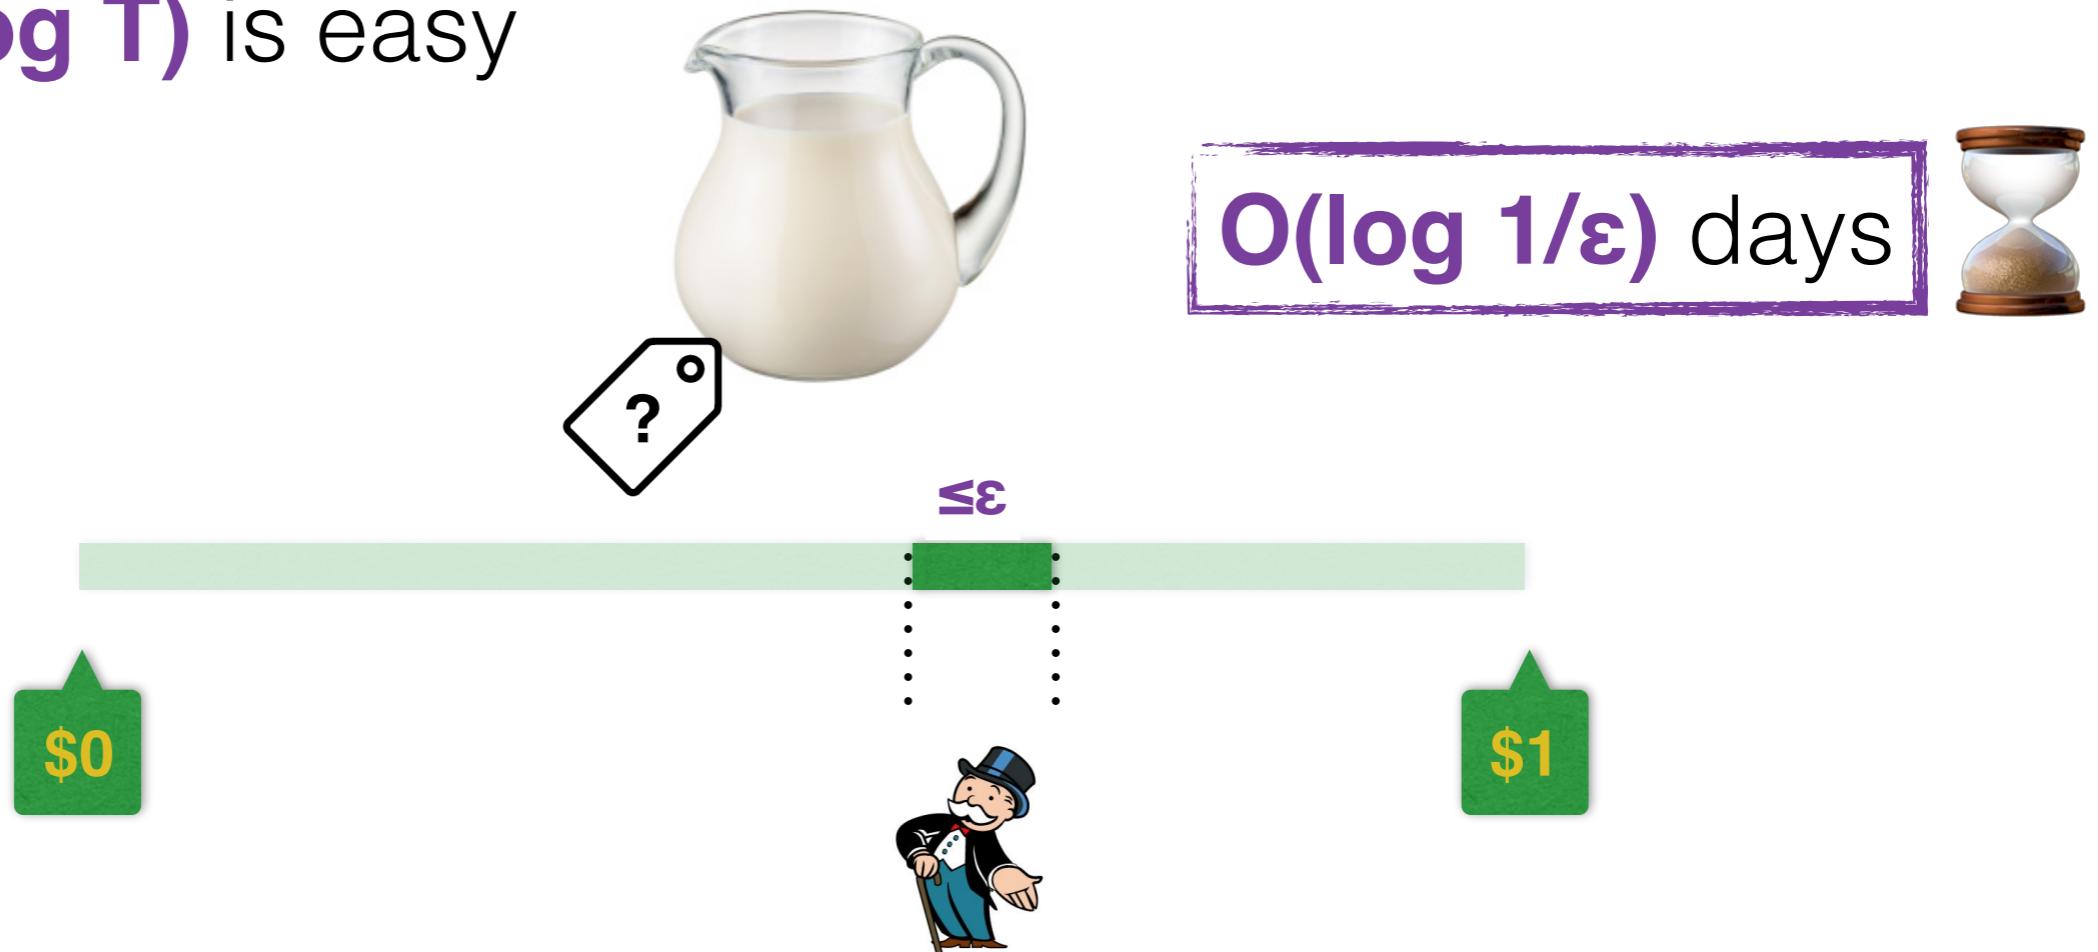
- Warm-up: 1-dimensional case



Dynamic Pricing Problem

- Warm-up: 1-dimensional case

- $O(\log T)$ is easy



Dynamic Pricing Problem

- Warm-up: 1-dimensional case
 - $O(\log T)$ is easy

Dynamic Pricing Problem

- Warm-up: 1-dimensional case
 - $O(\log T)$ is easy
 - $O(\log \log T)$ is a cute puzzle [Kleinberg, Leighton '03]

Dynamic Pricing Problem

- Warm-up: 1-dimensional case
 - $O(\log T)$ is easy
 - $O(\log \log T)$ is a cute puzzle [**Kleinberg, Leighton '03**]
- Back to d-dimensional case

Dynamic Pricing Problem

- Warm-up: 1-dimensional case
 - $O(\log T)$ is easy
 - $O(\log \log T)$ is a cute puzzle [Kleinberg, Leighton '03]
- Back to d-dimensional case
 - $O(d^2 \log T)$ [Cohen, Lobel, Paes Leme '16]

Dynamic Pricing Problem

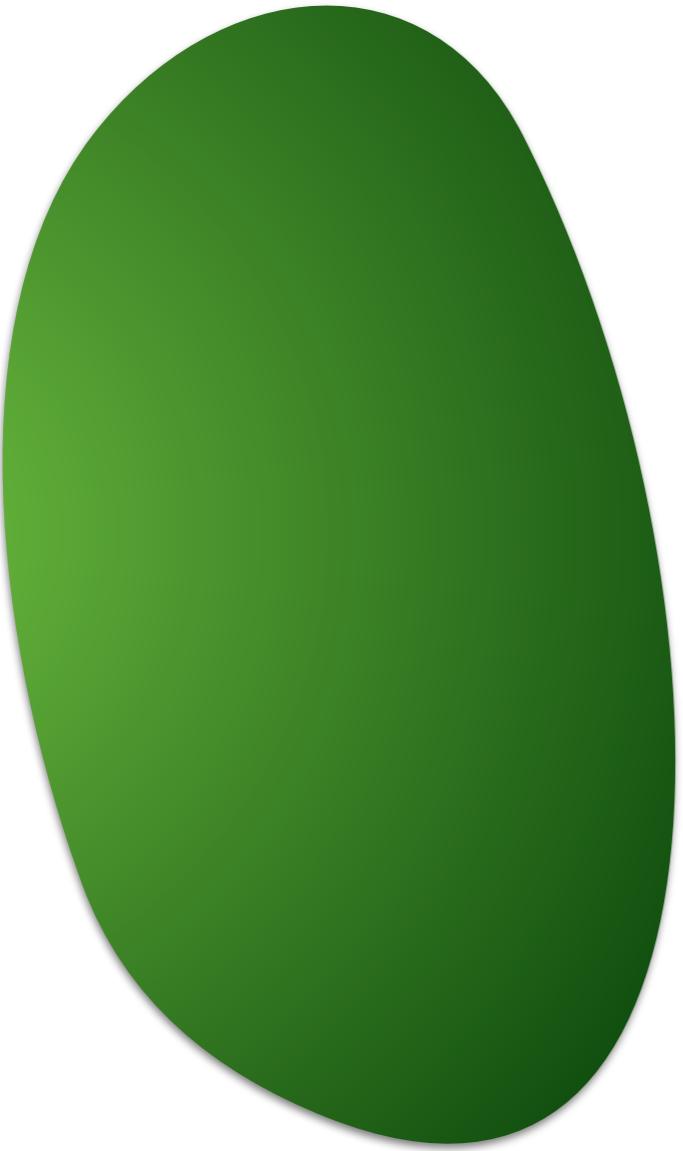
- Warm-up: 1-dimensional case
 - $O(\log T)$ is easy
 - $O(\log \log T)$ is a cute puzzle [Kleinberg, Leighton '03]
- Back to d-dimensional case
 - $O(d^2 \log T)$ [Cohen, Lobel, Paes Leme '16]
 - $\Omega(d \log \log T)$ lower bound

Dynamic Pricing Problem

- Warm-up: 1-dimensional case
 - $O(\log T)$ is easy
 - $O(\log \log T)$ is a cute puzzle [Kleinberg, Leighton '03]
- Back to d-dimensional case
 - $O(d^2 \log T)$ [Cohen, Lobel, Paes Leme '16]
 - $\Omega(d \log \log T)$ lower bound
 - Can we bridge the gap?

Dynamic Pricing Problem

- Warm-up: 1-dimensional case
 - $O(\log T)$ is easy
 - $O(\log \log T)$ is a cute puzzle [Kleinberg, Leighton '03]
- Back to d-dimensional case
 - $O(d^2 \log T)$ [Cohen, Lobel, Paes Leme '16]
 - $\Omega(d \log \log T)$ lower bound
 - $O(d \log T)$ [this work]

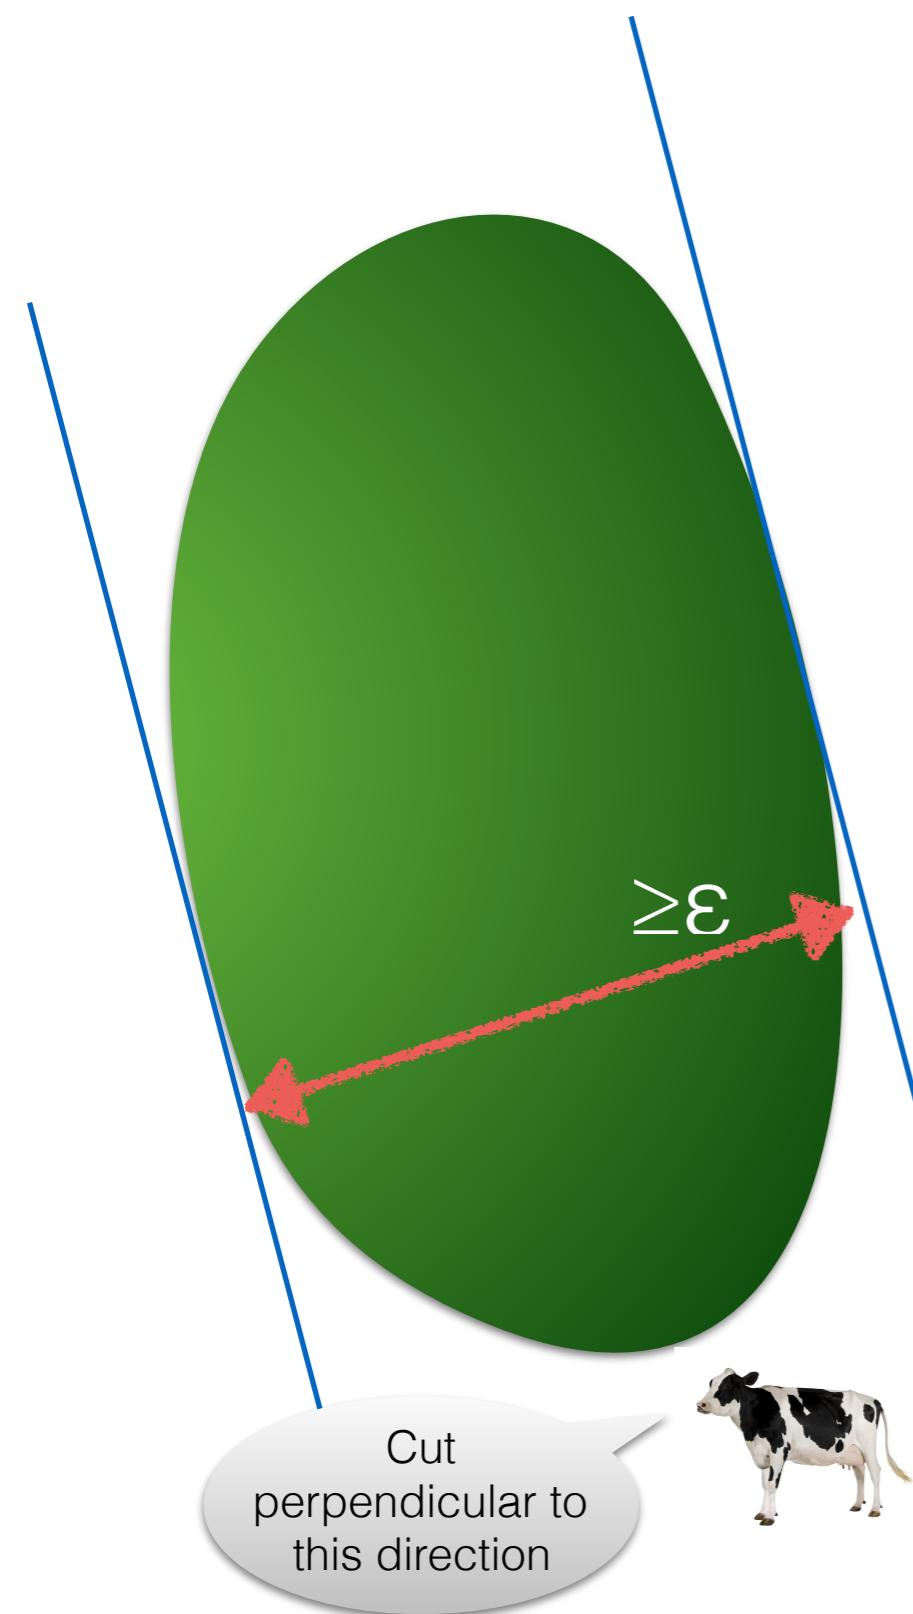


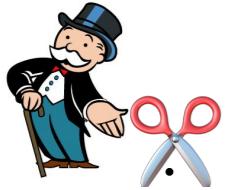
Model

- start with d -dimensional convex body (possible prices for items)

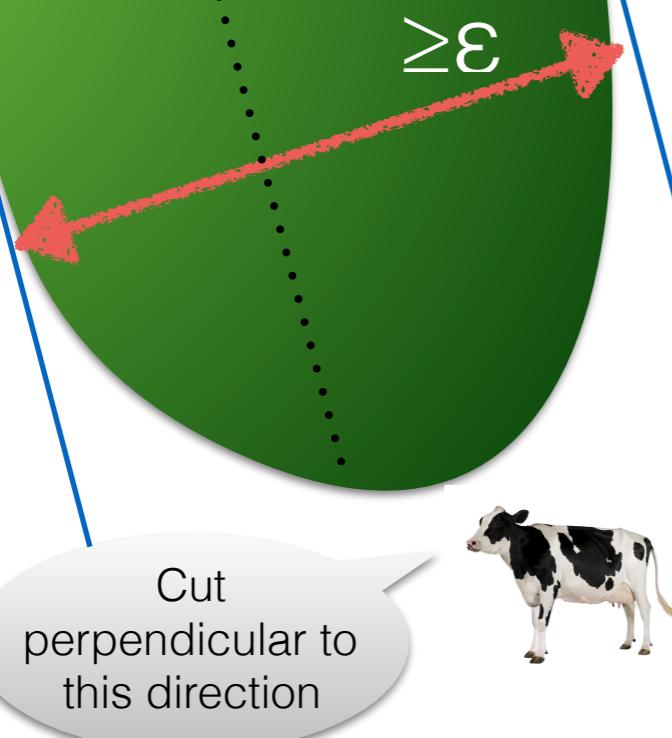
Model

- start with d -dimensional convex body (possible prices for items)
- customer gives vector of item quantities u , $\|u\| = 1$, with the promise that $w(K, u) \geq \epsilon$, asks for price

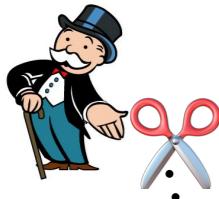




Model



- start with d -dimensional convex body (possible prices for items)
- customer gives vector of item quantities u , $\|u\| = 1$, with the promise that $w(K, u) \geq \epsilon$, asks for price
- player chooses price c , and partitions $K = K_1 \cup K_2$, where $K_1 = \{x \in K : x^\top u \leq c\}$ and $K_2 = K \setminus K_1$

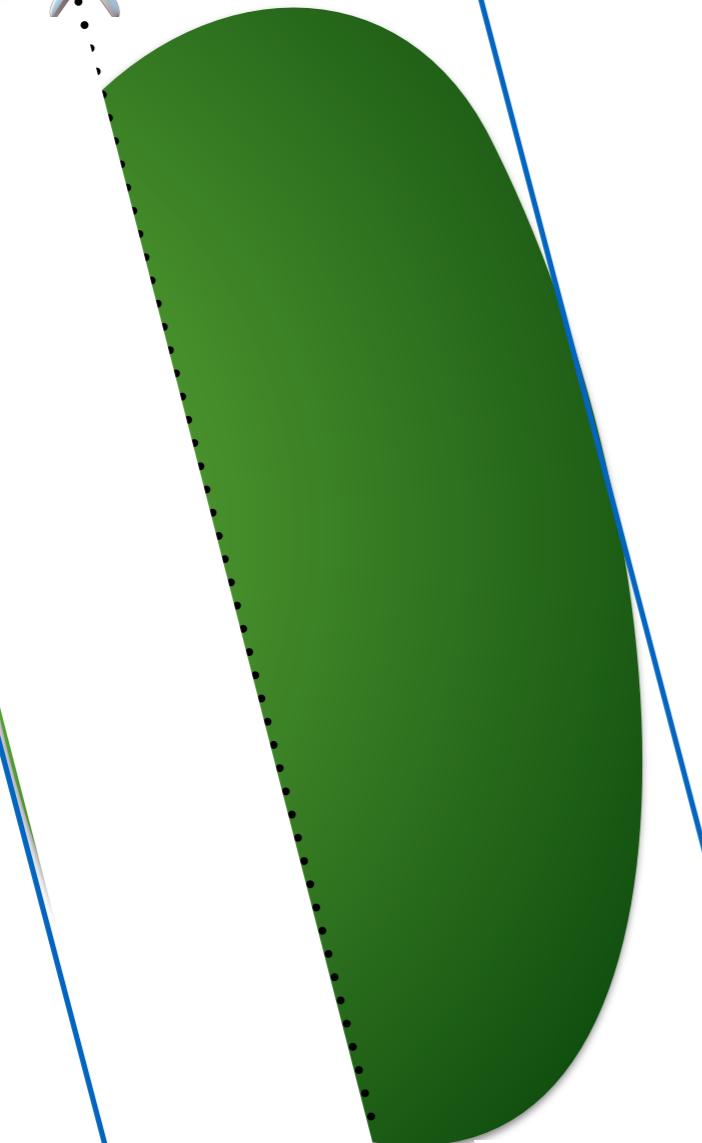
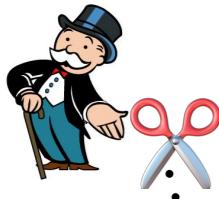


Model

Keep this half



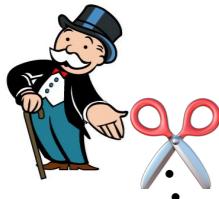
- start with d -dimensional convex body (possible prices for items)
- customer gives vector of item quantities u , $\|u\| = 1$, with the promise that $w(K, u) \geq \epsilon$, asks for price
- player chooses price c , and partitions $K = K_1 \cup K_2$, where $K_1 = \{x \in K : x^\top u \leq c\}$ and $K_2 = K \setminus K_1$
- customer picks which half contains the true prices, update K accordingly



Model

- start with d -dimensional convex body (possible prices for items)
- customer gives vector of item quantities u , $\|u\| = 1$, with the promise that $w(K, u) \geq \epsilon$, asks for price
- player chooses price c , and partitions $K = K_1 \cup K_2$, where $K_1 = \{x \in K : x^\top u \leq c\}$ and $K_2 = K \setminus K_1$
- customer picks which half contains the true prices, update K accordingly
- repeat until no valid directions remaining





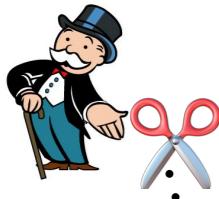
Model

Keep this half



- start with d -dimensional convex body (possible prices for items)
- customer gives vector of item quantities \mathbf{u} , $\|\mathbf{u}\| = 1$, with the promise that $w(K, \mathbf{u}) \geq \epsilon$, asks for price
- player chooses price c , and partitions $K = K_1 \cup K_2$, where $K_1 = \{x \in K : x^\top \mathbf{u} \leq c\}$ and $K_2 = K \setminus K_1$
- customer picks which half contains the true prices, update K accordingly
- repeat until no valid directions remaining

Theorem: there exists a strategy for cutting such that $w(K, \mathbf{u}) \leq \epsilon$ for all \mathbf{u} after $O(d \log (d/\epsilon))$ iterations.



Model

Keep this half

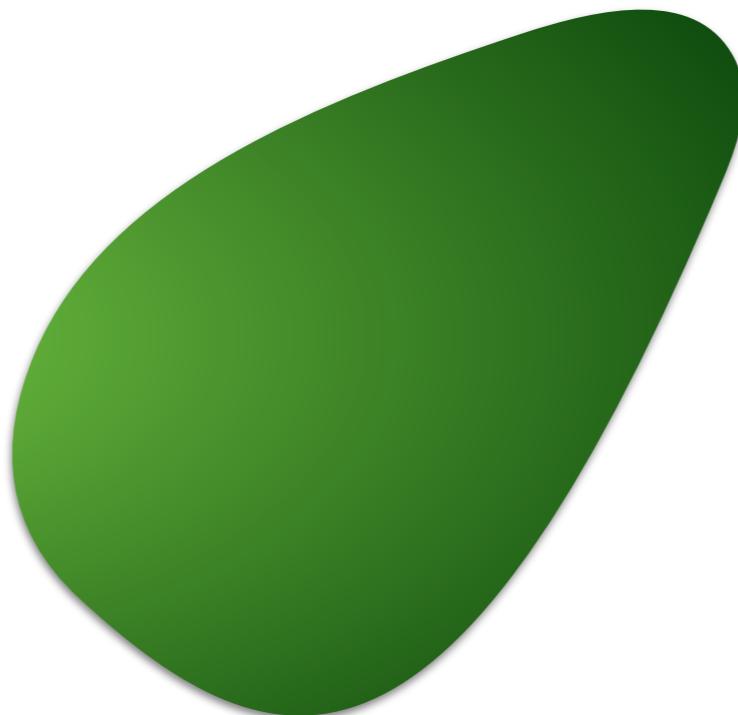


- start with d -dimensional convex body (possible prices for items)
- customer gives vector of item quantities u , $\|u\| = 1$, with the promise that $w(K, u) \geq \epsilon$, asks for price
- player chooses price c , and partitions $K = K_1 \cup K_2$, where $K_1 = \{x \in K : x^\top u \leq c\}$ and $K_2 = K \setminus K_1$
- customer picks which half contains the true prices, update K accordingly
- repeat until no valid directions remaining

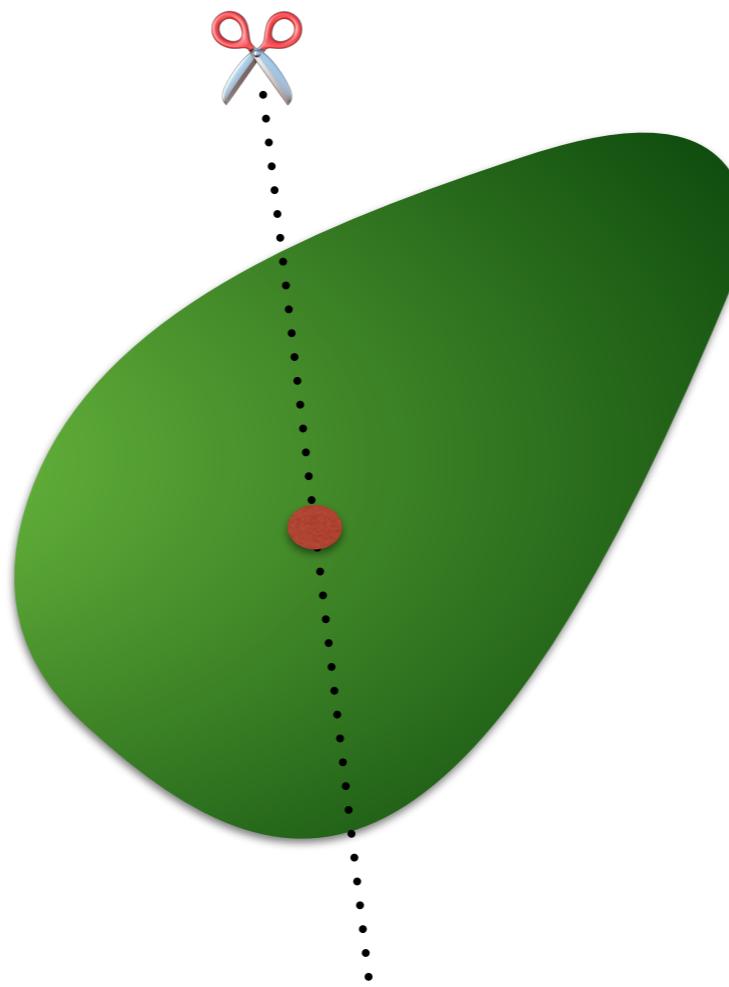
Theorem: there exists a strategy for cutting such that $w(K, u) \leq \epsilon$ for all u after $O(d \log (d/\epsilon))$ iterations.

Set $\epsilon = d \log T / T$, regret is at most $T\epsilon + O(d \log (d/\epsilon)) = O(d \log T)$.

Convex Geometry Toolbox

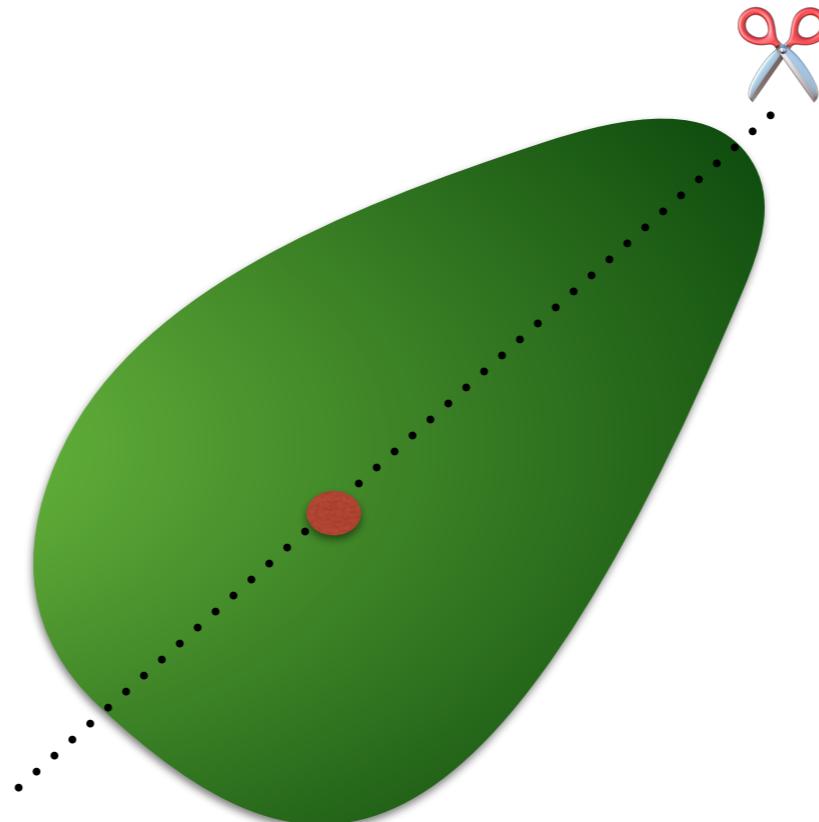


Convex Geometry Toolbox



Grunbaum's Theorem: any cut through the centroid partitions K in two pieces of roughly equal volume (largest/smallest $\leq e$).

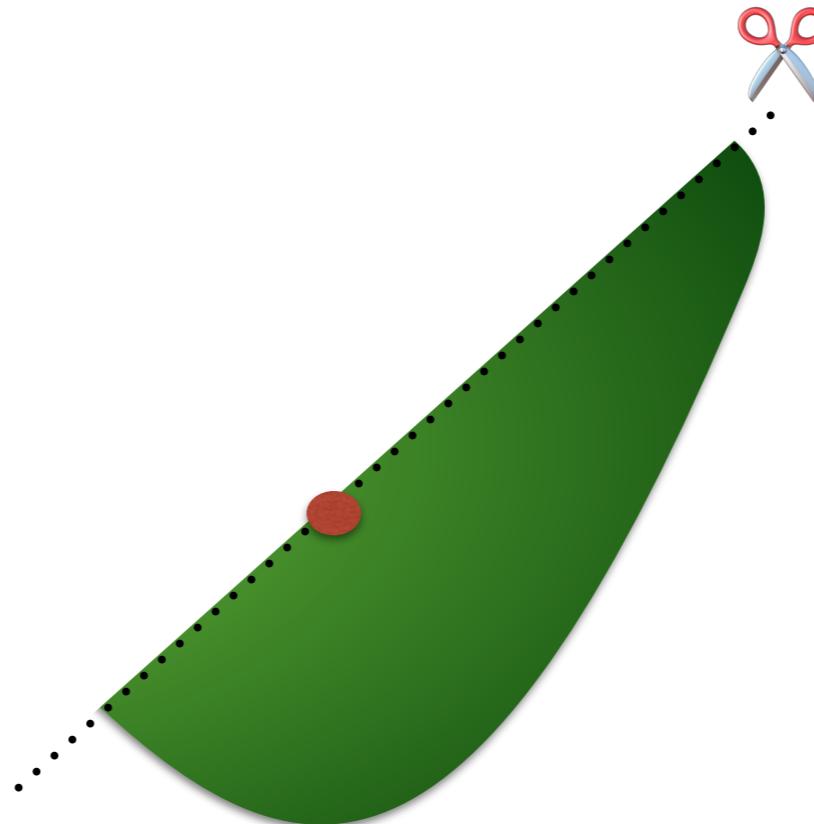
Convex Geometry Toolbox



Not quite
sufficient

Grunbaum's Theorem: any cut through the centroid partitions \mathbf{K} in two pieces of roughly equal volume (largest/smallest $\leq e$).

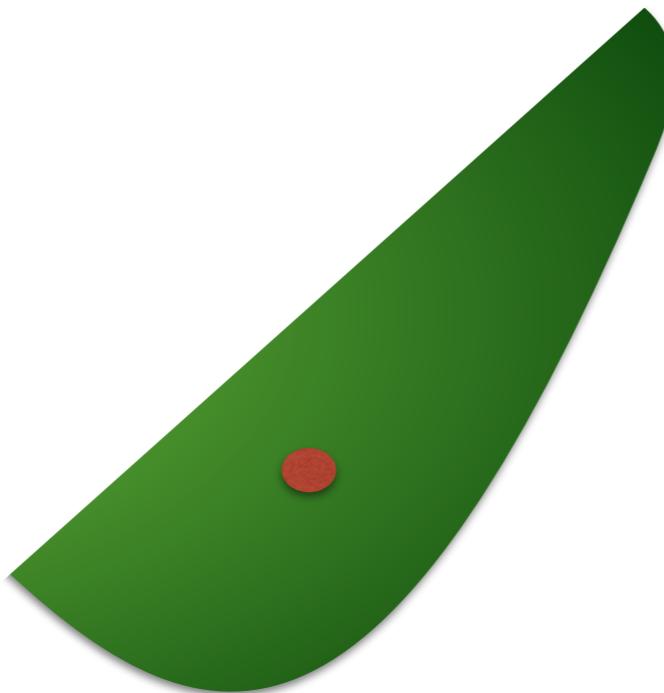
Convex Geometry Toolbox



Not quite
sufficient

Grunbaum's Theorem: any cut through the centroid partitions \mathbf{K} in two pieces of roughly equal volume (largest/smallest $\leq \epsilon$).

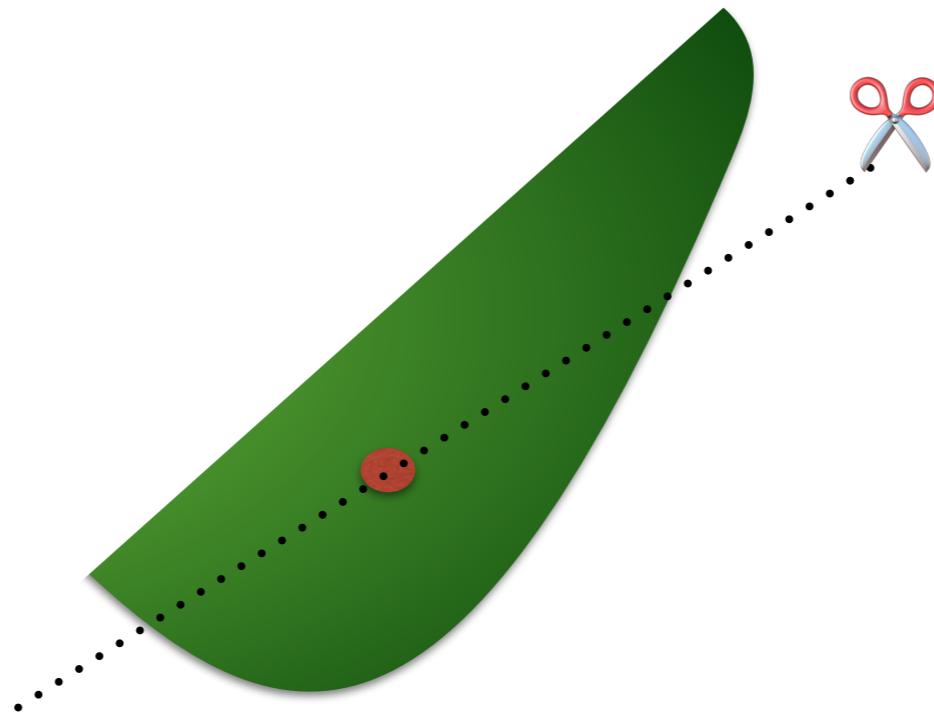
Convex Geometry Toolbox



Not quite
sufficient

Grunbaum's Theorem: any cut through the centroid partitions \mathbf{K} in two pieces of roughly equal volume (largest/smallest $\leq e$).

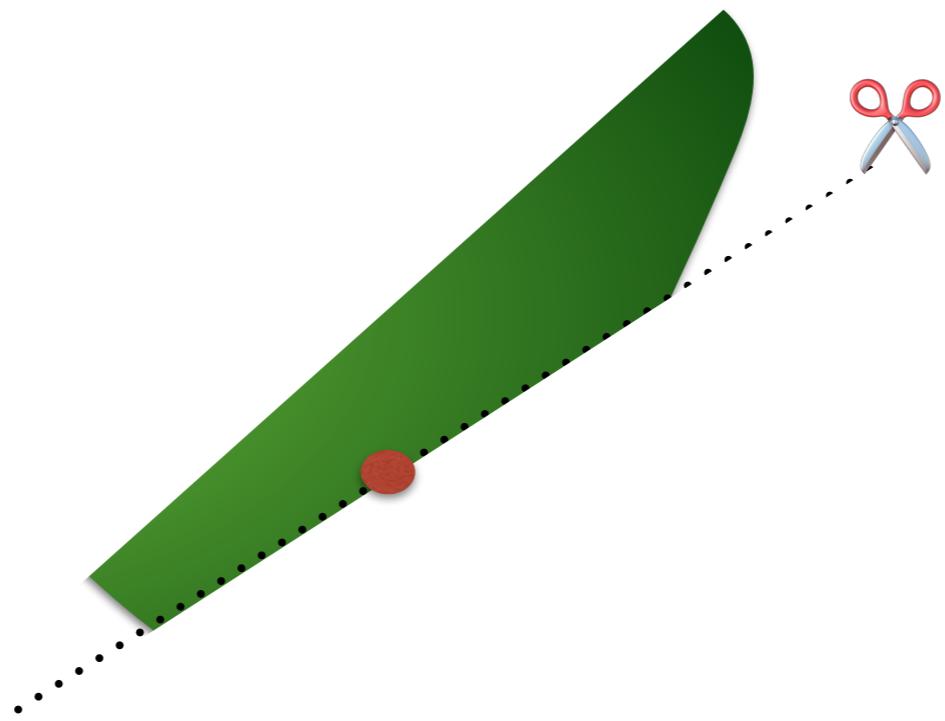
Convex Geometry Toolbox



Not quite
sufficient

Grunbaum's Theorem: any cut through the centroid partitions \mathbf{K} in two pieces of roughly equal volume (largest/smallest $\leq e$).

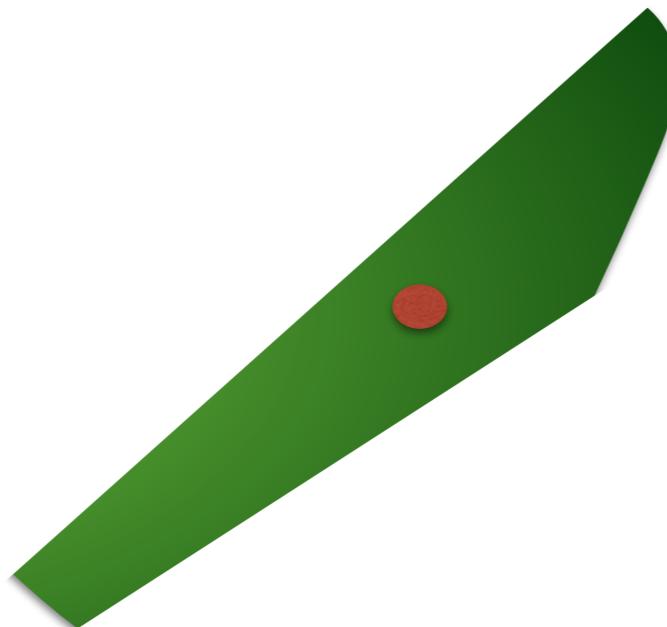
Convex Geometry Toolbox



Not quite
sufficient

Grunbaum's Theorem: any cut through the centroid partitions \mathbf{K} in two pieces of roughly equal volume (largest/smallest $\leq e$).

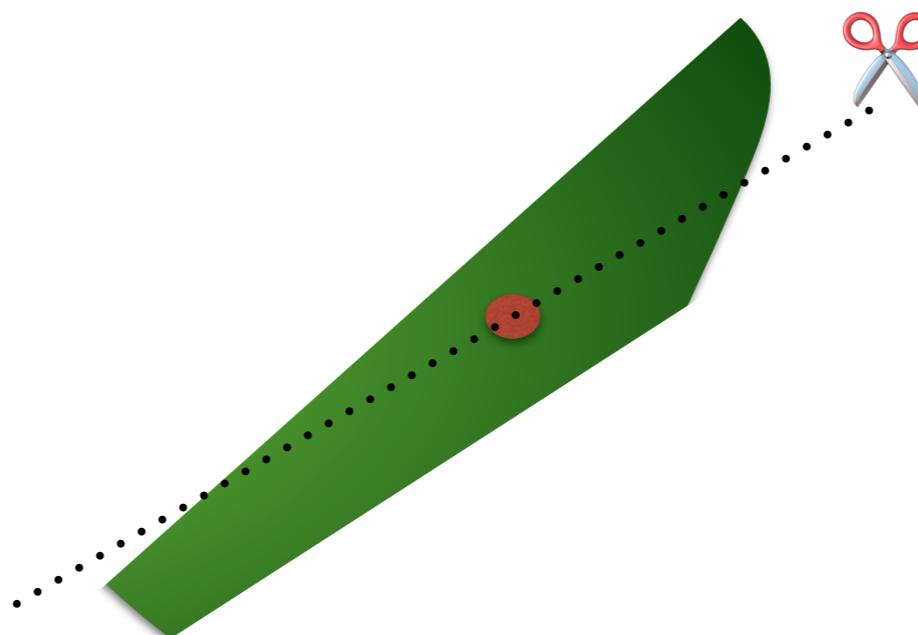
Convex Geometry Toolbox



Not quite
sufficient

Grunbaum's Theorem: any cut through the centroid partitions K in two pieces of roughly equal volume (largest/smallest $\leq e$).

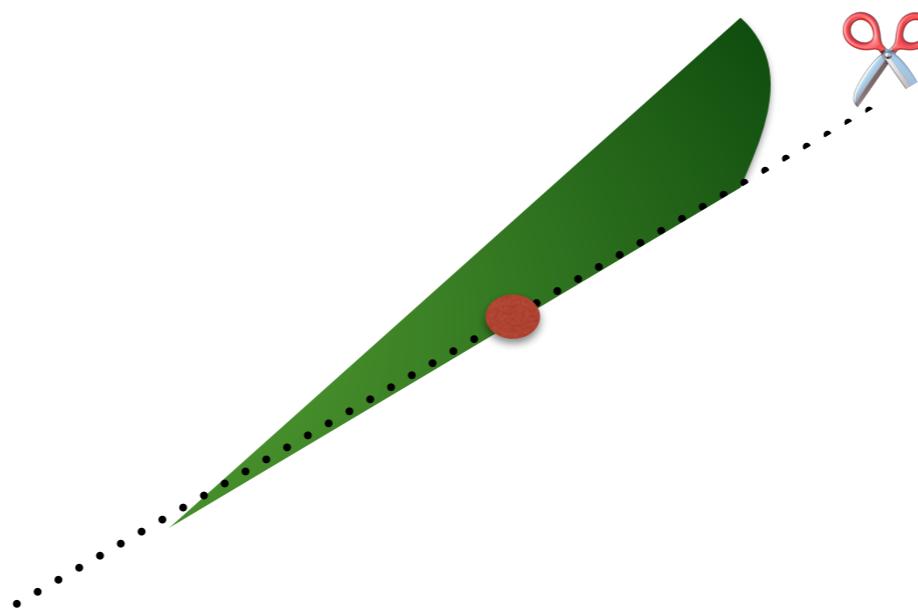
Convex Geometry Toolbox



Not quite
sufficient

Grunbaum's Theorem: any cut through the centroid partitions \mathbf{K} in two pieces of roughly equal volume (largest/smallest $\leq e$).

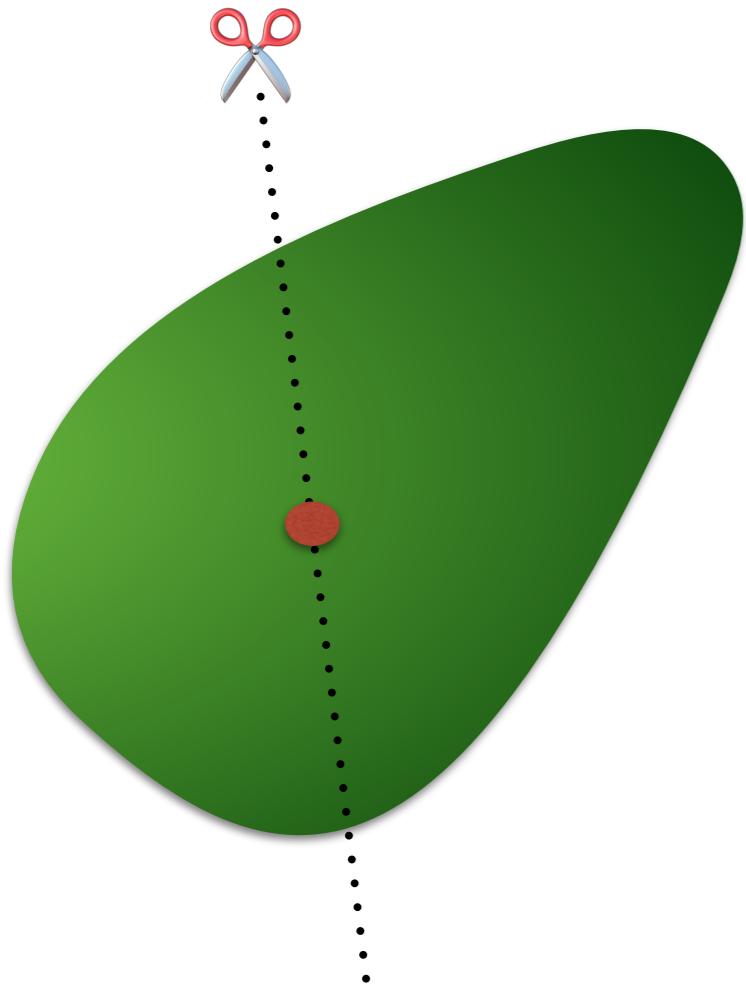
Convex Geometry Toolbox



Not quite
sufficient

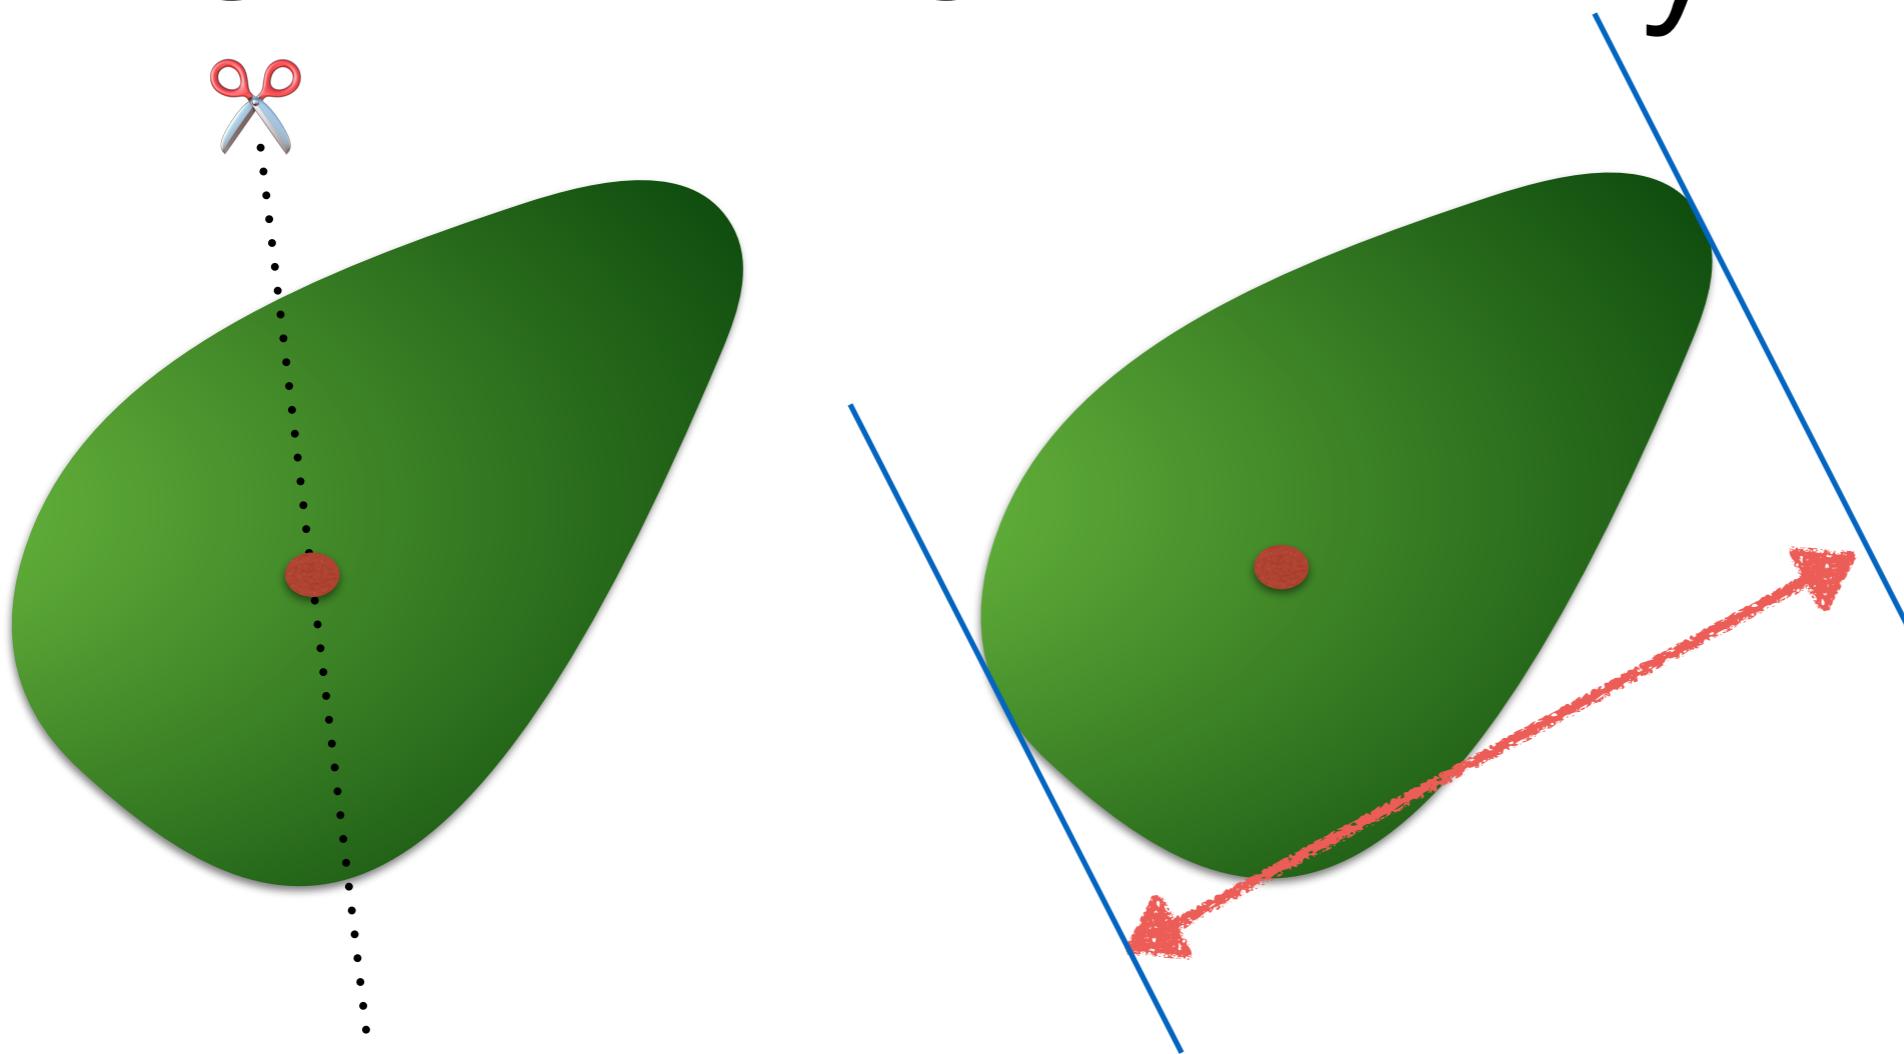
Grunbaum's Theorem: any cut through the centroid partitions \mathbf{K} in two pieces of roughly equal volume (largest/smallest $\leq e$).

Convex Geometry Toolbox



Grunbaum's Theorem: Any cut through $c(K)$ partitions K in two pieces with volume ratio at most e .

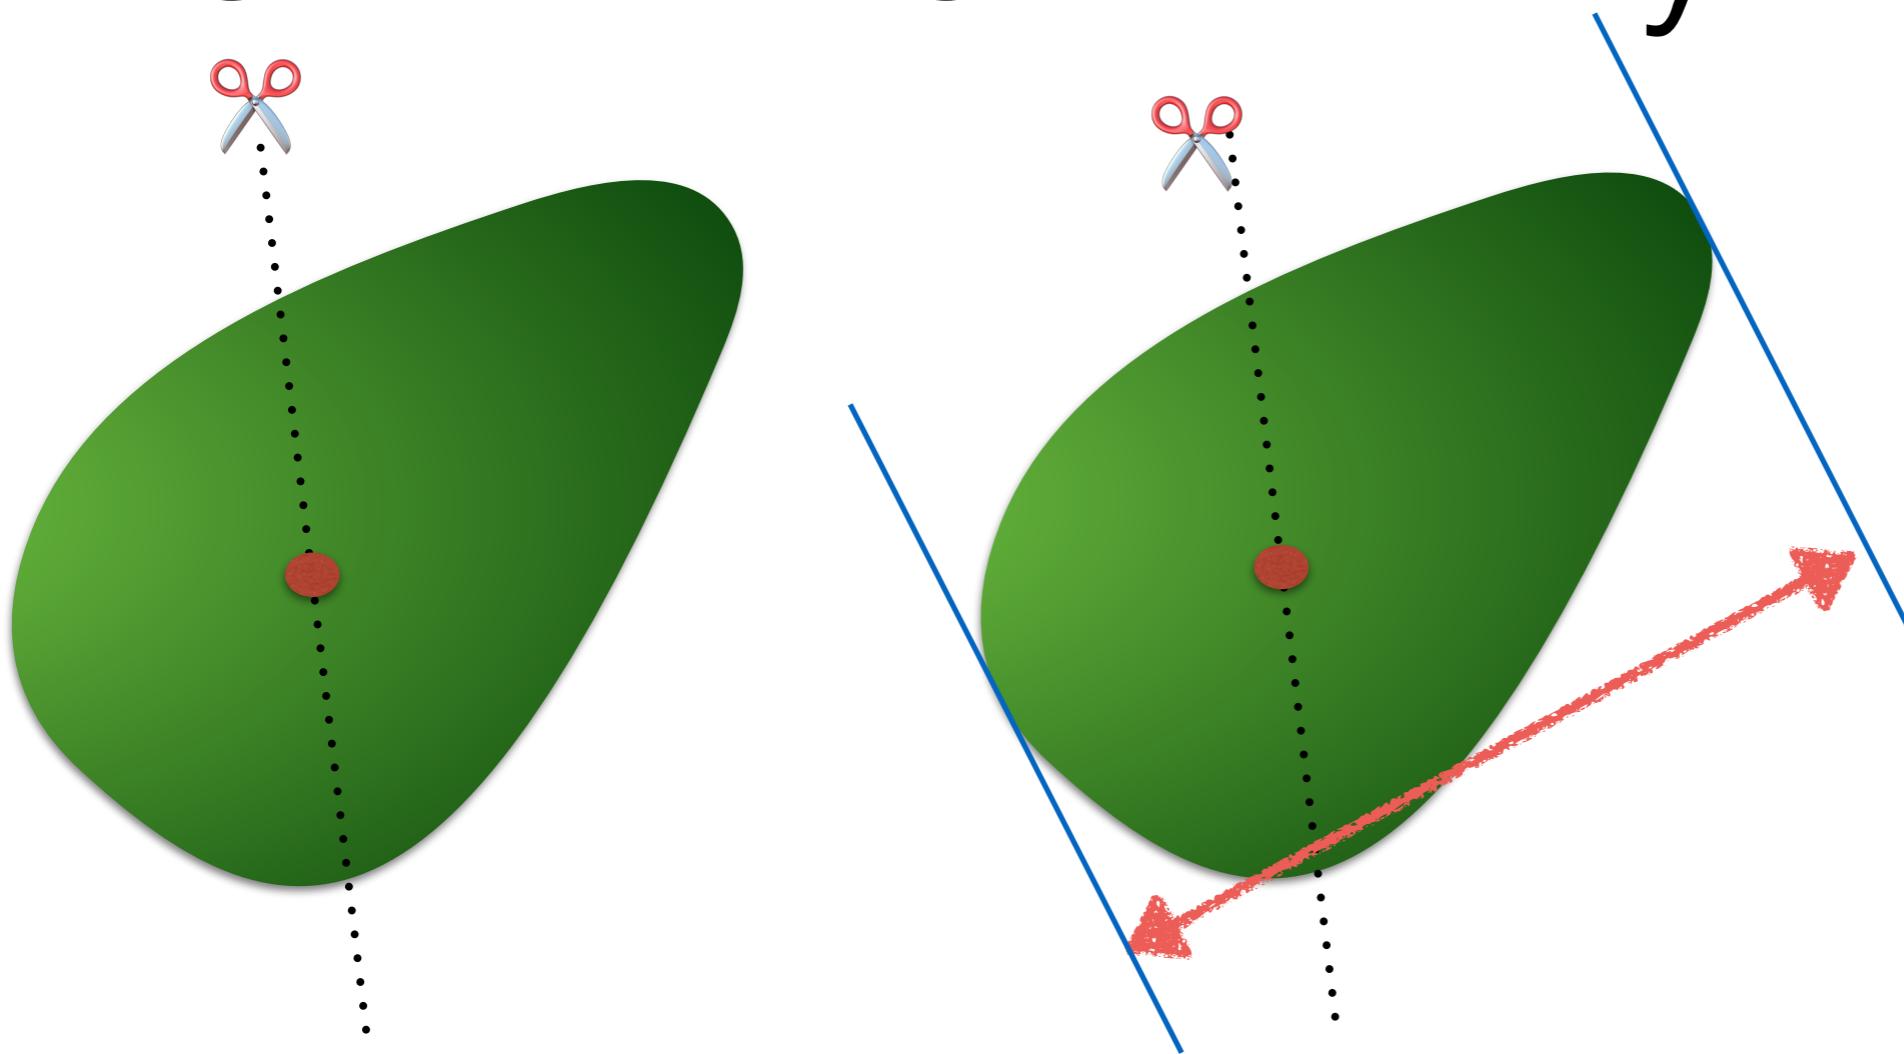
Convex Geometry Toolbox



Grunbaum's Theorem: Any cut through $c(K)$ partitions K in two pieces with volume ratio at most e .

Directional Grunbaum (this work): A cut through $c(K)$ partitions $K=K_1 \cup K_2$ such that $w(K_i) \geq w(K,u) / (d+1)$, for all u .

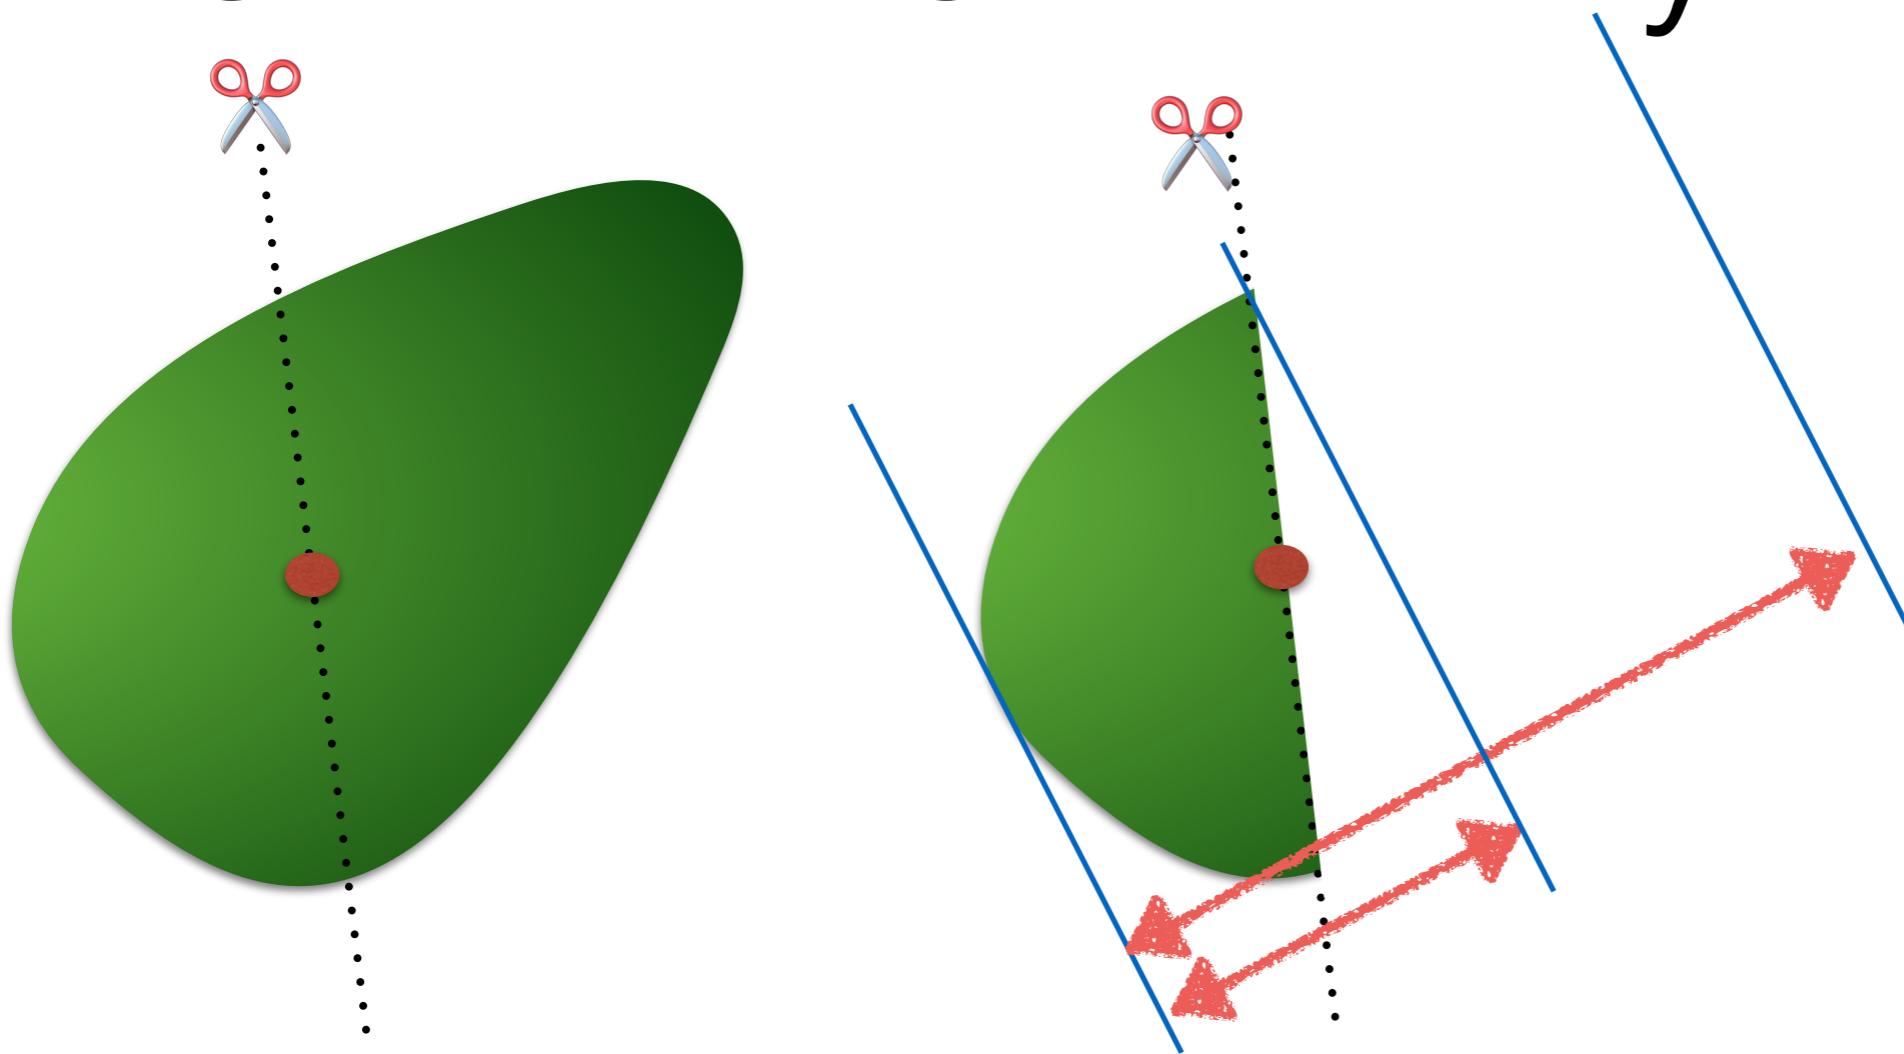
Convex Geometry Toolbox



Grunbaum's Theorem: Any cut through $c(K)$ partitions K in two pieces with volume ratio at most e .

Directional Grunbaum (this work): A cut through $c(K)$ partitions $K=K_1 \cup K_2$ such that $w(K_i) \geq w(K,u) / (d+1)$, for all u .

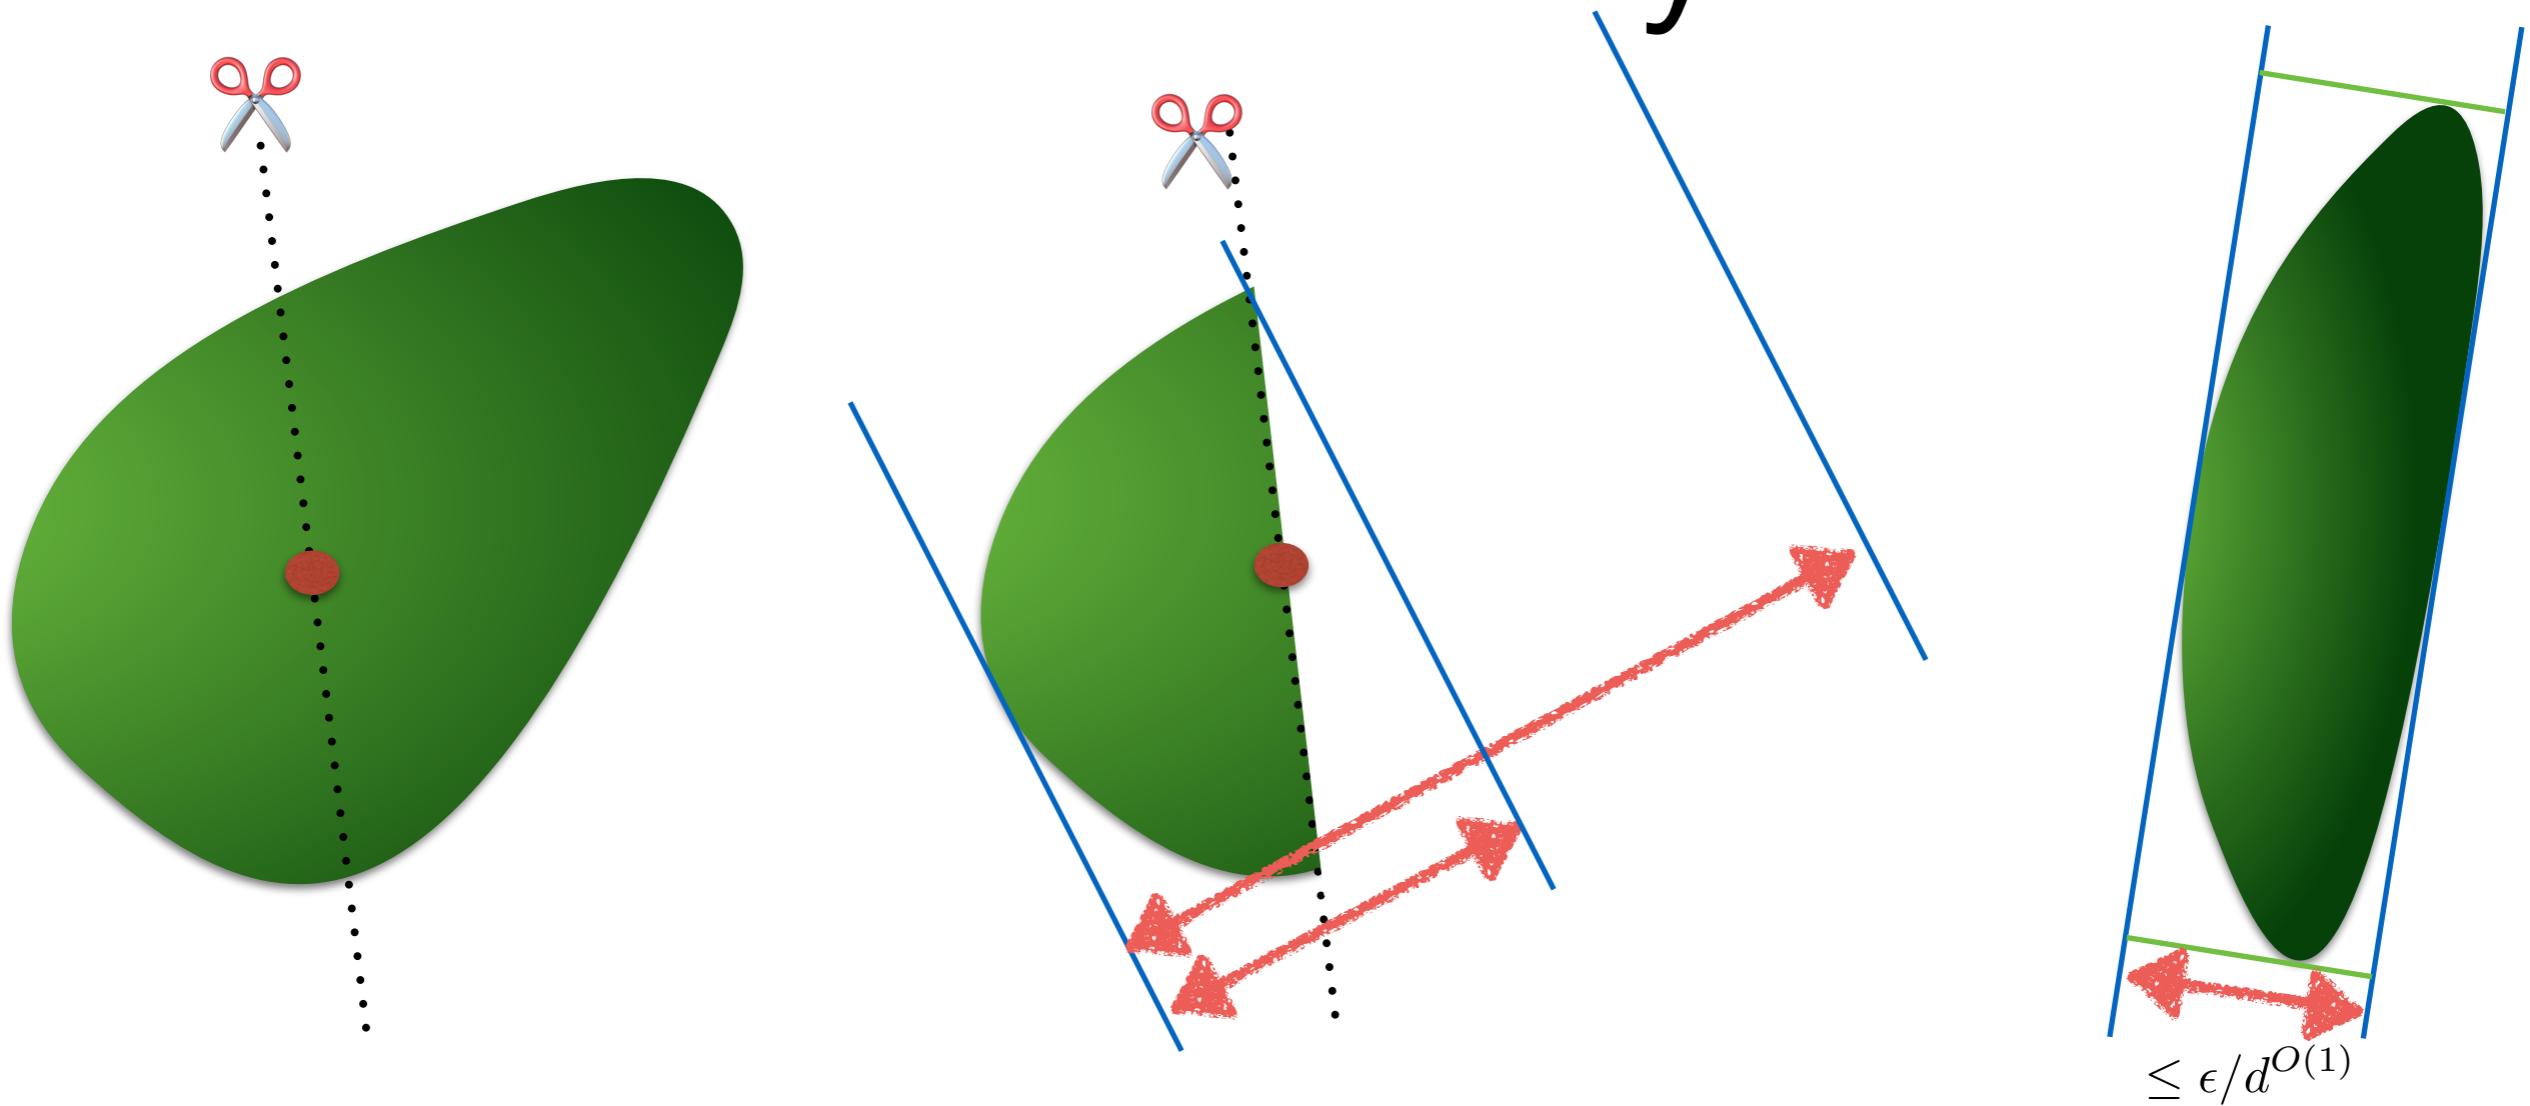
Convex Geometry Toolbox



Grunbaum's Theorem: Any cut through $c(K)$ partitions K in two pieces with volume ratio at most e .

Directional Grunbaum (this work): A cut through $c(K)$ partitions $K = K_1 \cup K_2$ such that $w(K_i) \geq w(K, u) / (d+1)$, for all u .

Convex Geometry Toolbox



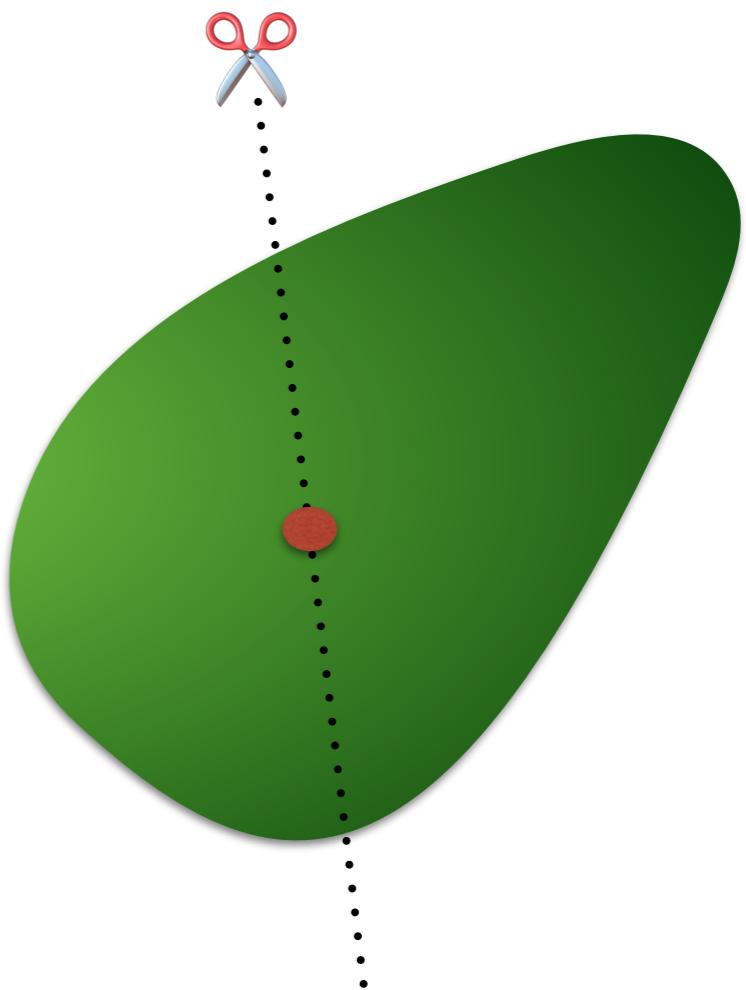
Grunbaum's Theorem: Any cut through $c(K)$ partitions K in two pieces with volume ratio at most e .

Directional Grunbaum (this work): A cut through $c(K)$ partitions $K = K_1 \cup K_2$ such that $w(K_i) \geq w(K, u) / (d+1)$, for all u .

Cylindrification (this work): If $w(K, u) \geq \delta$ for all u , and S is $(d-1)$ -dimensional, then $\text{Vol}(\text{Projs } K) \leq (d+1)^2/\delta \text{ Vol}(K)$.

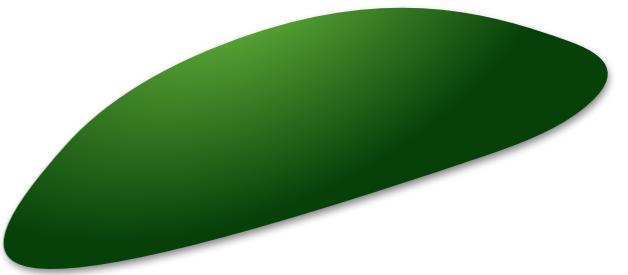
Algorithm

1. Cut through centroid.



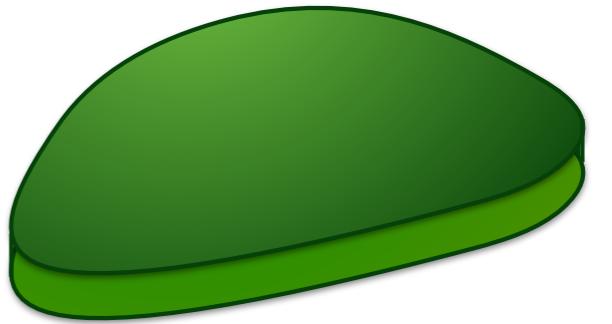
Algorithm

1. Cut through centroid.
2. When $w(K, u) \leq \delta = \epsilon/d^{O(1)}$, cylindrify.



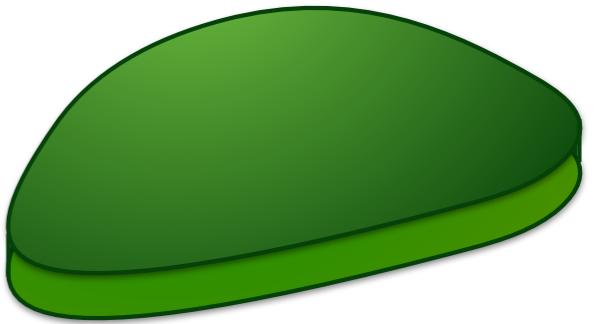
Algorithm

1. Cut through centroid.
2. When $w(K, u) \leq \delta = \epsilon/d^{O(1)}$, cylindrify.



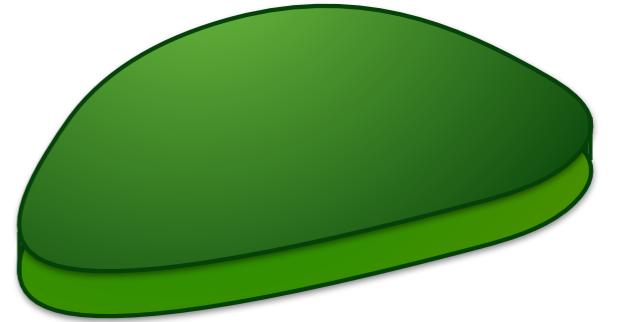
Algorithm

1. Cut through centroid.
2. When $w(K, u) \leq \delta = \epsilon/d^{O(1)}$, cylindrify.
3. Small volume blow-up from $\leq d$ cylindrifications, overall need to reduce volume by a factor of $(d/\epsilon)^{O(d)}$.



Algorithm

1. Cut through centroid.
2. When $w(K, u) \leq \delta = \epsilon/d^{O(1)}$, cylindrify.
3. Small volume blow-up from $\leq d$ cylindrifications, overall need to reduce volume by a factor of $(d/\epsilon)^{O(d)}$.

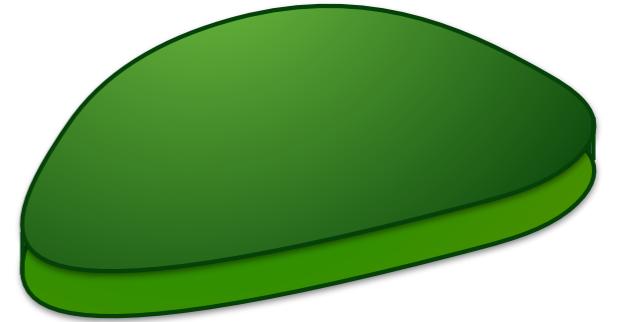


POLYNOMIAL TIME ALGORITHM

1. Compute approximate centroid, requires almost uniform sampling [LS '93]

Algorithm

1. Cut through centroid.
2. When $w(K, u) \leq \delta = \epsilon/d^{O(1)}$, cylindrify.
3. Small volume blow-up from $\leq d$ cylindrifications, overall need to reduce volume by a factor of $(d/\epsilon)^{O(d)}$.

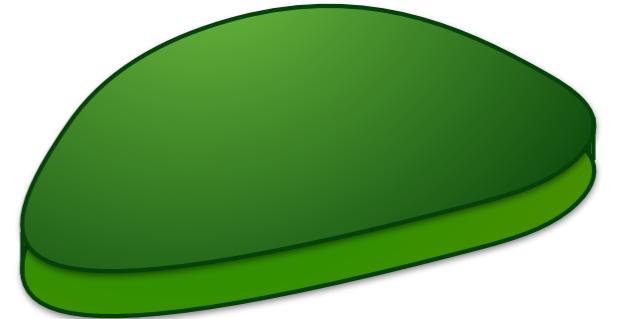


POLYNOMIAL TIME ALGORITHM

1. Compute approximate centroid, requires almost uniform sampling [LS '93]
2. Robust versions of our theorems

Algorithm

1. Cut through centroid.
2. When $w(K, u) \leq \delta = \epsilon/d^{O(1)}$, cylindrify.
3. Small volume blow-up from $\leq d$ cylindrifications, overall need to reduce volume by a factor of $(d/\epsilon)^{O(d)}$.



POLYNOMIAL TIME ALGORITHM

1. Compute approximate centroid, requires almost uniform sampling [LS '93]
2. Robust versions of our theorems

Open: obtain regret $O(d \log d + d \log \log T)$ by combining with the 1-dimensional trick.

Thank you!

